

# L01: Introduction to Database Systems

DSAN-6300 & PPOL-6810

**Databases & SQL Programming**

Irina Vayndiner

August 28 & 31, 2023



GEORGETOWN UNIVERSITY

# Outline for today's class

## Welcome to DSAN-6300 / PPOL-6810

- Introductions
- Course Overview
- Introduction to Database Systems
  - What is a Database System
  - Application and Purpose of Database Systems
  - History of Database Systems
  - Data Models
  - Database Languages
  - Database Architecture

# Introducing your teacher



Irina Vayndiner  
MITRE

Email: [iv95@georgetown.edu](mailto:iv95@georgetown.edu)

Interests:

- Databases, Big Data
- Storytelling
- Improvisational theater (long-form)

Other courses I taught/teach at Georgetown:

- Massive Data Fundamentals (ANLY-502, HIDS-511, PPOL-567)
- Storytelling with Data (ANLY-599)

# Introducing TAs

## TAs

- Raghav Sharma
- Ramdayal Rewaria (Lead DSAN)
- Yanyan Li
- Parsa Keyvani
- Shenghao Wang
- Pablo Jimenez
- Peijin Li (Lead PPOL)

**TA Office Hours & locations:** published on Canvas

**STUDENTS POLL**

# Overall Course Topics

- Introduction to Relational Models and Schema Design
- SQL Programming (from beginners to advanced!)
- How databases work
  - Query Processing & Optimization methods
  - Concurrency and Transactions
  - Database Architectures
- Data storage including column-oriented and distributed storage
- Intro to SQL Big Data Concepts and Processing
- NoSQL and other non-relational databases

- Hands-on software and tools
  - MySQL Workbench on your laptop
  - RDS in AWS Cloud
  - MongoDB in AWS Cloud

# Class Logistics

- **Classroom**
  - Mon: CBN #204
  - Tue: St Marys #110
- **Lectures & Labs Day/Time**
  - **Mon** for ANLY-6300-01: 3:30-6pm EST
  - **Thu** for ANLY-6300-02 and PPOL-6810: 12:30-3pm EST
- There will be some exceptions
  - Check Canvas announcements regularly!



# Class Logistics (continued)

- Typically, each class will involve:
  - Lecture
  - Break
  - Lab hands-on work (based on the Lecture material)
- Communication outside of class:
  - Use Discussion boards on Canvas if ran into an issue or have a question
  - Use email for administrative issues
    - Need to inform
      - NLT 10am the day of the class if not able to attend class or attending on Zoom
      - NLT 3pm on the day assignment is due
    - When you email me, please cc at least one TA
      - Ram for DSAN, Peijin for PPOL
      - Unless there is an issue that is not appropriate for TAs
- Please plan to use your laptops for every class

# Class Materials on Canvas

We will use Canvas for the Class materials

- Announcements
  - Check frequently, or sign up for email alerts
- Syllabus
- Modules, including dates of tests
  - Data of the last test might shift!
- Lecture Slides (published after the lecture!)
- Assignments, e.g.
  - Usually will publish a reminder when an Assignment is published
- Deadlines
- Grades

**Important: all Sections of this Course on Canvas will look like one section:  
ANLY-6300-01: You are still in the section you originally signed for!**

# Class Policies

- Class participation in classroom is expected, it is part of the grade
  - Ask & answer questions
- Assignments are due **on**
  - **HWs & Quizzes: Tuesdays at 11:59pm**
  - Usually around 10-14 days for HWs, one week for Quizzes
- Labs are due after the class on
  - **Thursday 11:59pm for Mon class**
  - **Tuesday 11:59pm for Thu class**
- **No late** submission for **Quizzes or Project**
- Canvas Locks at the due time!
- All assignments are individual assignments
- Follow the Georgetown Student Pledge
  - Will be painful if you do not!

# Class Late Submission Policy

- Late Submission Policy
  - All non-test deliverables can be up to **24 hours late with NO PENALTY.**
  - Any assignment submitted after that grace period will result in **reduction of points** as follows:
    - HW Assignments: -5% per day for up to 5 days
    - All other Assignments: -10% per day for up to 5 days
  - The Midterm and Final text **CANNOT** be rescheduled
    - **Midterm: Thu 10/19 and Mon 10/23**
    - **Final test on zoom: Tue 12/5 at 10:30am for all 3 sections**
    - Please make a note of these dates NOW

# DSAN-6300 / PPOL 6810, by the numbers

Years taught:	2019	2020	2021	2022	2023
Class sessions:	13	14	13	14	14
Class length:	2hrs 30min	2hrs 30min	2hr 30min	2h 30min	2h 30min
Enrolled students:	26	72	92	95	96
Num of Sections	1	2	2	3	3
Homework assignments:	3	3	3	3	3
Mini-Project:	1	1	1	1	1
Online Quizzes:	4	4	4	4	4
Labs	0	9	9	9	9
In-Class Tests	2 (Closed books)	2 (Open books)	2 (Open books)	2 (Closed & Open Books)	2 (Closed and Open books)

# Tentative Class Dates, check Canvas for updates!

	DSAN-6300-01 Mon	DSAN-6300-02& PPOL-6810 Thu
1	8/28	8/31
2	9/5 (Tue!)	9/7
3	9/11	9/14
4	9/18	9/21
5	9/25 (recording)	9/28
6	10/2	10/5
7	10/16	10/12
8 (midterm)	10/23	10/19
9	10/30	10/26
10	11/6	11/2
11	11/13	11/9
12	11/20	11/16
13	11/27	11/30
14 (test)	12/5 (Tue, 10:30am-1pm)	12/5 (Tue, 10:30-1pm)

## No classes

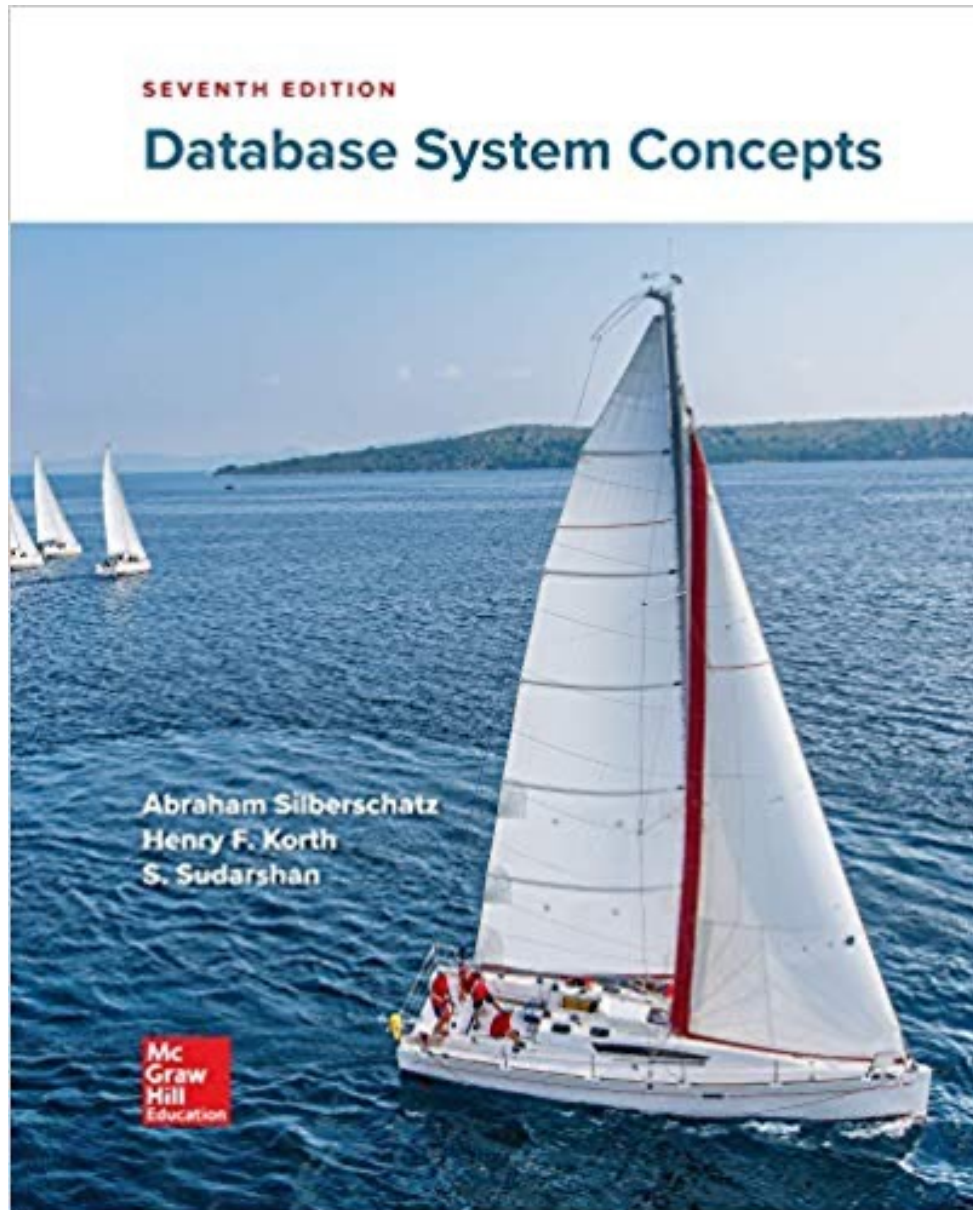
Mon, 9/4 Labor Day

Mon, 10/9: Mid Semester  
Holiday

Mon, 9/25: Yom Kippur

Class added: Mon 12/4 for  
in lieu of 8/24 (all 3 sections)

# Textbook



- **7<sup>th</sup> Edition(2020)**
- 6<sup>th</sup> edition is outdated!
- Will use this book's materials for many of the Course lectures

# Class Deliverables: What you need to do!

Course Activity	Weight
Class participation including in-class Labs	12%
Four Homeworks, including one Project	40%
Four On-line Quizzes	8%
Two In-Class Tests (midterm 15% and “final” 25%)	40%



# Grading

- Final point count is 100%, the letter grade will follow standard guidelines for Fall 2023

## Grades

- A:  $\geq 92.5$
- A-: 89.5 - 92.49
- B+: 87.99 - 89.49
- B: 81.5 - 87.98
- B-: 79.5 - 81.49

# Grading Rubric for Homework Assignments & Project

- We will look at the results files and the scripts.
- If the result files are exactly what is expected, in the proper format, etc., we may run your scripts to make sure they produce the output. If everything works, you will get full credit for the problem.
- If the results files are not what is expected, we will look and run your code and provide partial credit wherever possible and applicable.
- Points **will** be deducted for each/any the following reasons:
  - Instructions are not followed (please read assignments carefully!)
  - Output is not in expected format (e.g. not sorted, missing fields, wrong delimiter, unusual characters in the files, etc.)
  - There are more files delivered than need to be
  - There are additional lines in the results files (whether empty or not)
  - Files are not the requested filename
  - The Assignment is late: see late submission policy

# What you need to take this course

- Textbook
- Commitment to Assignments
  - Quiz will take up to 1hr
  - Finishing the Lab will take about 2hrs on avg
  - **HW will take more time**— plan ahead!
- Use of Amazon Educate Cloud (we will explain!)
  - We will use Amazon Web Services (AWS) for hands-on work
- Hardware:
  - Your laptop (OS: Mac or Windows w/ Cygwin)
    - If you have OS Windows on your laptop, please email TA Paolo
- Software (Installation Instructions will follow)
  - MySQL Workbench
  - RDS in AWS Cloud
  - MongoDB in AWS Cloud

# Tentative Schedule for Fall 2023: Monitor Canvas for updates

Class num	Class	Labs (due <b>Tue</b> on the week after the class)	Book Chapters	Homeworks	Quizzes	Due on <b>Tue</b> of the week
1	L1. Introduction, Database Systems		1			
2	L2. Database Design Using the E-R Model	Lab: Setup	6		Q1	
3	L3: Relational Model and Relational Algebra	Lab: Setup	2, 7	HW1		Q1
4	L4 SQL 1	Lab: SQL 1	3			
5	L5. SQL 2	Lab: SQL 2	3, 4	HW2	Q2	HW1
6	L6. SQL 3	Lab: SQL 3	4,5			Q2
7	L7. Advanced Topics in SQL and Database Design	Lab: Adv Topics	15, 16	HW3		HW2
8	Midterm					H2
9	L9. Intro to Big Data	Lab: TBD	10, 11, 12, 14		Q3	HW3
10	L10. Non-relational DB + Guest Speaker	Lab: MongoDB		HW4 (Project)		
11	L11. Data Warehousing, Data Analytics, and Transactions	Lab: Analytics	TBD			Q4
12	L12. Database Optimization	Lab: Db Optimization	11, 15		Q4	Q3
13	L13. Data Management for the Enterprise					
14	Test					Project

# Why you need to take this course



“As I looked around at some open data science positions, I think about 99% of them required a working knowledge of SQL“

# Introduction to Database Systems



# Databases are everywhere

- A **database** is a collection of data, managed over a period of time
  - Big: Google web crawls, sales records at Walmart
  - Small: a personal Christmas card list, a recipe file
  - Specialized: medical images, Facebook links, a hybrid auto design
  - You are in Georgetown student database
- Databases are increasingly found in unexpected places
  - Are your photos in Apple cloud?
  - There is a personal autocomplete data on your phone or email
  - Sometimes people are not aware they use a database since accessing database forms is part of everyday life

The screenshot shows a travel booking application window titled "Paradise Holidays - [Booking Details]". The interface includes a menu bar (File, Booking, View, Reports, Tools, Help) and a toolbar with buttons for Main, Avail, Agent, Requests, Trans, Pax, Receipts, and Cancel. The form is divided into several sections:

- Booking Details:** Includes fields for Booking Number (1000), Booking Name (Vincent), Other Code, Date Created (20/11/01), and Date Modified.
- Travel Details:** Includes fields for Destination (ON JETTY), Repeat (No), Initial Source (GCM JETTY), and Status Price (BROCHURE PRICE).
- Customer Details:** Includes fields for Green Name (Vincent), Designation (No), Address (4219), State (QLD), Country (AUSTRALIA), Home Phone (3816 2885), Work Phone, Mobile, Fax, and Email.
- Booking Summary:** Includes fields for Price (1475), Disc (BROCHURE P), GST Free (0), Type (Confirmed), NP (1475), and Due Date (Mon 20/5/02).
- Payment Details:** Includes fields for Service (P2), Status (Confirmed), Amount (975.00), and Outboard to Di (1 Complete).

At the bottom, there is a table showing the booking history:

Date	Assistant	Status	Complete	Due Date
Wed 28/11/01	Emily	1 Complete	12440	
Th 30/11/01	Letika	1 Complete	12489	
Sat 20/1/02	Dinny	1 Complete	14069	

# Why are Databases Important to Study

- Lots of benefits from understanding data principles
  - Databases are crucial components of complex software systems
    - How does Walmart decide what to put on shelves?
    - How Amazon decides what to recommend to you?
    - How does 7-11 decide what to put on the shelves next to each other?
  - Database perspective will allow you to understand better how the whole system works
  - The principles themselves may even be your biggest takeaway from this course...

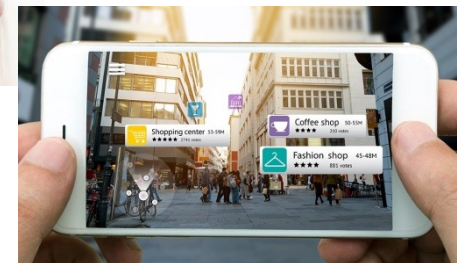
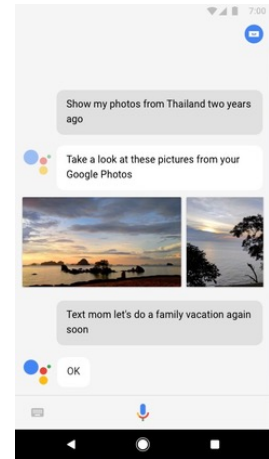


# Some Database Examples

- Databases touch all aspects of our lives
- Enterprise Information
  - Sales: customers, products, purchases
  - Accounting: payments, receipts, assets
  - Human Resources: Information about employees, salaries, payroll taxes.
- Manufacturing: management of production, inventory, orders, supply chain.
- Banking and finance
  - Customer information, accounts, loans, and banking transactions.
  - Credit card transactions
  - Finance: sales and purchases of financial instruments (e.g., stocks and bonds; storing real-time market data
- Universities: registration, grades, etc.

# Examples of Data in Everyday Life

- Facebook, LinkedIn, Tweeter, SnapChat, etc.
- Targeted advertising
- Google Assistant (“Hey Google”), Siri, Alexa
  - “Gets you just the right information at just the right time”
- Wearables:
  - Smart Glasses: wearable computer with a head-mounted display
    - Google Glass (one of the first):
      - Why didn't it succeed?
      - Resurfaced as Glass Enterprise Edition
    - Currently many types of Smartglasses
  - Apple Watch, etc.
  - Health monitoring wearables
- Netflix
  - Uses Data to predict if TV show will be a hit with audiences before it is produced
- Virtual reality
  - Data researchers are betting on VR technologies to explore new ways to visualize and analyze complex and dynamic datasets.



# Examples of Government Data

- Sensor Data
  - Biosecurity
  - Emergency preparedness
  - Logistics, e.g. for optimized decisions about logistics
  - Economic Data, e.g. unemployment rate
  - Financial Data, e.g. compliance & fraud detection
  - Climate modeling
  - Environmental Data, e.g. environmental monitoring & mitigation
  - Cybersecurity
  - Healthcare, e.g. public health, epidemics early detection
- 
- Big Data Research is supported by Government, for example:
    - Department of Energy (DoE)
    - National Science Foundation (NSF)



# What is a DBMS

- We talked about databases. What is DBMS?
- A DBMS is a DataBase Management System
  - A software tool making it easier to manage databases
  - Over 50 years of research and engineering behind modern DBMSs
  - A huge business (billions of \$ per year)
  - In this course we will cover RDBMS (Relational DBMS), and also other DBMS's (nosql, semi-structured, key-value, etc)
- Q: Can we have a database without a DBMS?

# What is a DBMS (continued)

- A database system (DBMS) is a collection of interrelated data and a set of programs that allow users to access, manage and use these data.
- Major DBMS elements are:
  - Data models
    - A collection of tools for describing data
  - Data access
    - Structured Query Language (SQL)
    - Language, such as C/C++, Java or Python, with embedded SQL queries that access the data in the database
    - Application programs that are used to interact with the database in this fashion.
      - Important term: Application Programming Interface (API)

# Why use DBMS

- Let's say you have 1TB of Public Policy data
- You could just write an application that stores data in disk files
- Some issues you would need to address:
  - Speed: linear access can be slow, optimization and indexing can provide dramatic speedup
  - Query/report interface: you need to write a program for every question; is everybody who needs information a developer?
  - Schema: how does a new user know where to find “Criminal Justice Policy”?
  - Data integrity: if you delete a topic, how do you ensure all related items are deleted?
  - Security: how do you limit who can query sensitive information?
  - Transactions: you may need to code up protection for concurrent users “stepping on” each other, and also against system crashes
- That can be **a lot of work!**

# Challenges of Database Applications

In the early days, database applications were built directly on top of file systems, which led to

- Data redundancy and inconsistency
  - Data is stored in multiple file formats resulting in duplication of information in different files
- Difficulty in accessing data
  - Need to write a new program to carry out each new task
- Data isolation
  - Consistency and completeness of data retrieved
- Integrity problems
  - Integrity constraints (e.g., account balance  $> 0$ ) become “buried” in program code rather than being stated explicitly
  - Hard to add new constraints or change existing ones

# Challenges of Database Applications (continued)

- Atomicity of updates
  - Failures may leave database in an inconsistent state with partial updates carried out
    - Example: Transfer of funds from one account to another should either complete or not happen at all
- Concurrent access by multiple users
  - Concurrent access needed for performance
  - Uncontrolled concurrent accesses can lead to inconsistencies
    - Ex: Two people reading a balance (say \$100) and updating it by withdrawing money (say \$100 each) at the same time
- Security problems
  - Hard to provide user access to some, but not all, data

**Database systems offer solutions to all the above problems**



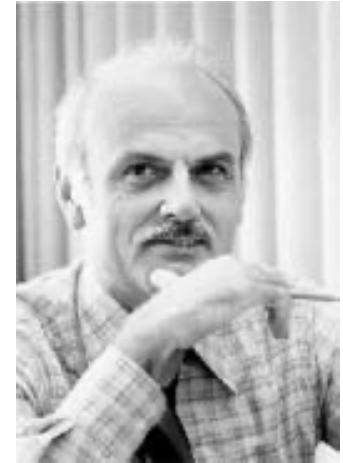
# Services that DBMS Offers

- A Standard Query Language (SQL)
- Automated query optimization, and sub-linear access through indexing
- Schematic data modeling / design
- Automated data integrity
- Security (access controls, audit logs, encryption, authentication ...)
- Crash-proof Transactions
- Concurrency control
- Massively parallel optimizations
- GUI's, APIs (e.g., JDBC), and a rich set of development tools
- Import/export: XML, JSON, CSV
- Sophisticated administration and tuning tools

If your application needs these, an DBMS is a powerful time-saving tool

# History of Database Systems

- 1832 — Semen Korsakov uses punch cards for data storage in Russia
- 1890 — Herman Hollerith used paper punch cards for US 1890 Census
  - 1896 Hollerith founded Tabulating Machine Company, which later becomes IBM
- 1950s and early 1960s
  - Data processing using magnetic tapes for storage; punch cards for input
  - 1964 — SABRE system goes online (American Airlines & IBM).
    - Two database models: network model (CODASYL) and hierarchical model (IMS)
- Late 1960s and 1970s
  - 1970-72 Edgar Frank "Ted" Codd defines the Relational Data Model
    - Codd won the ACM Turing Award for this work
    - IBM Research begins System R prototype
  - UC Berkeley (Michael Stonebraker) begins Ingres prototype
  - Oracle releases first commercial relational database
  - High-performance (for the era) transaction processing



INGRES

ORACLE®

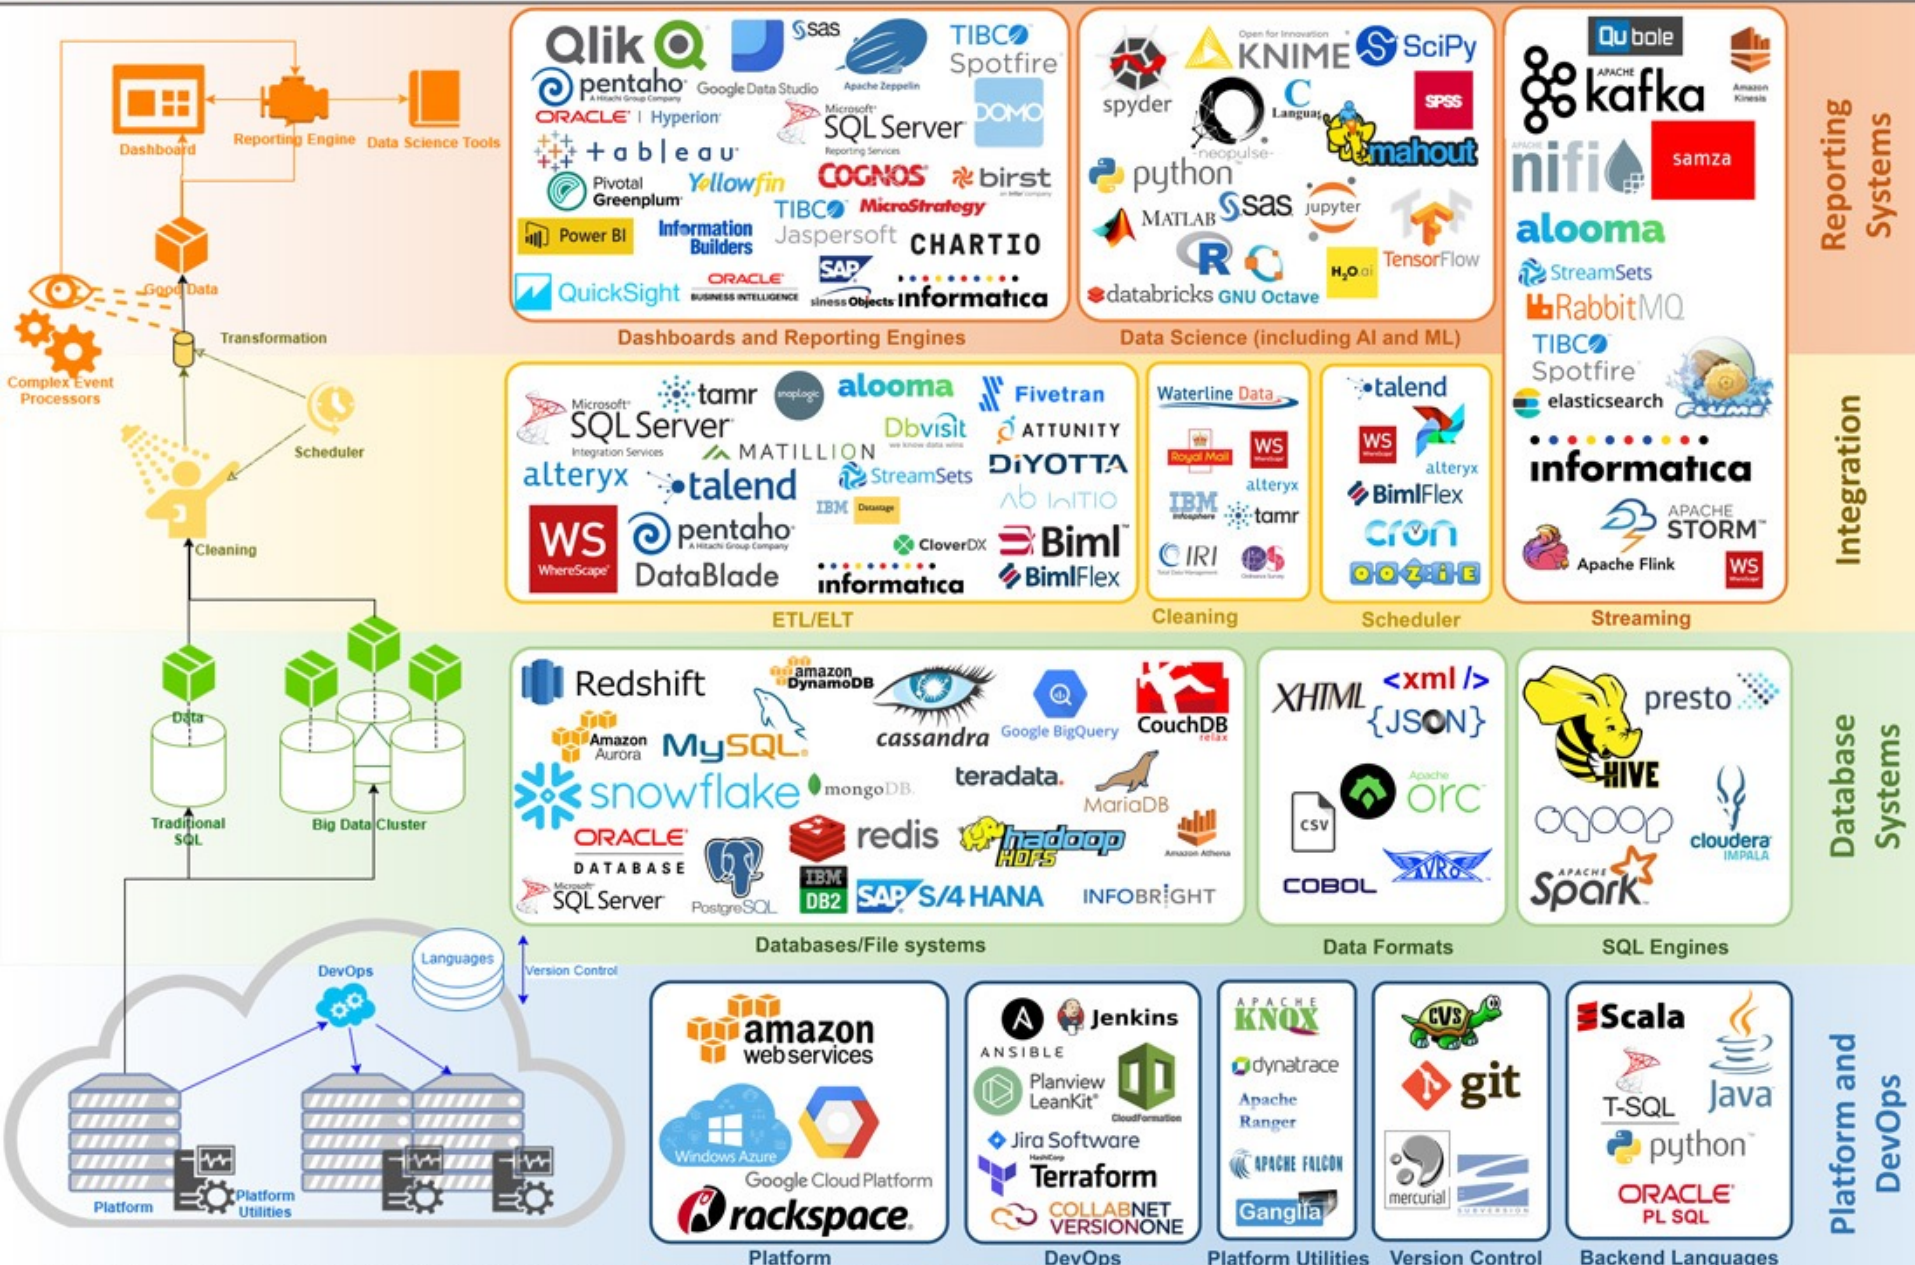
# History of Database Systems (Continued)

- 1980s
  - Research relational prototypes evolve into commercial systems
    - SQL becomes industrial standard
  - Parallel and distributed database systems
    - IBM, Teradata
  - Object-oriented database systems
- 1990s
  - Large decision support and data-mining applications
  - Large multi-terabyte data warehouses
  - Emergence of Web commerce



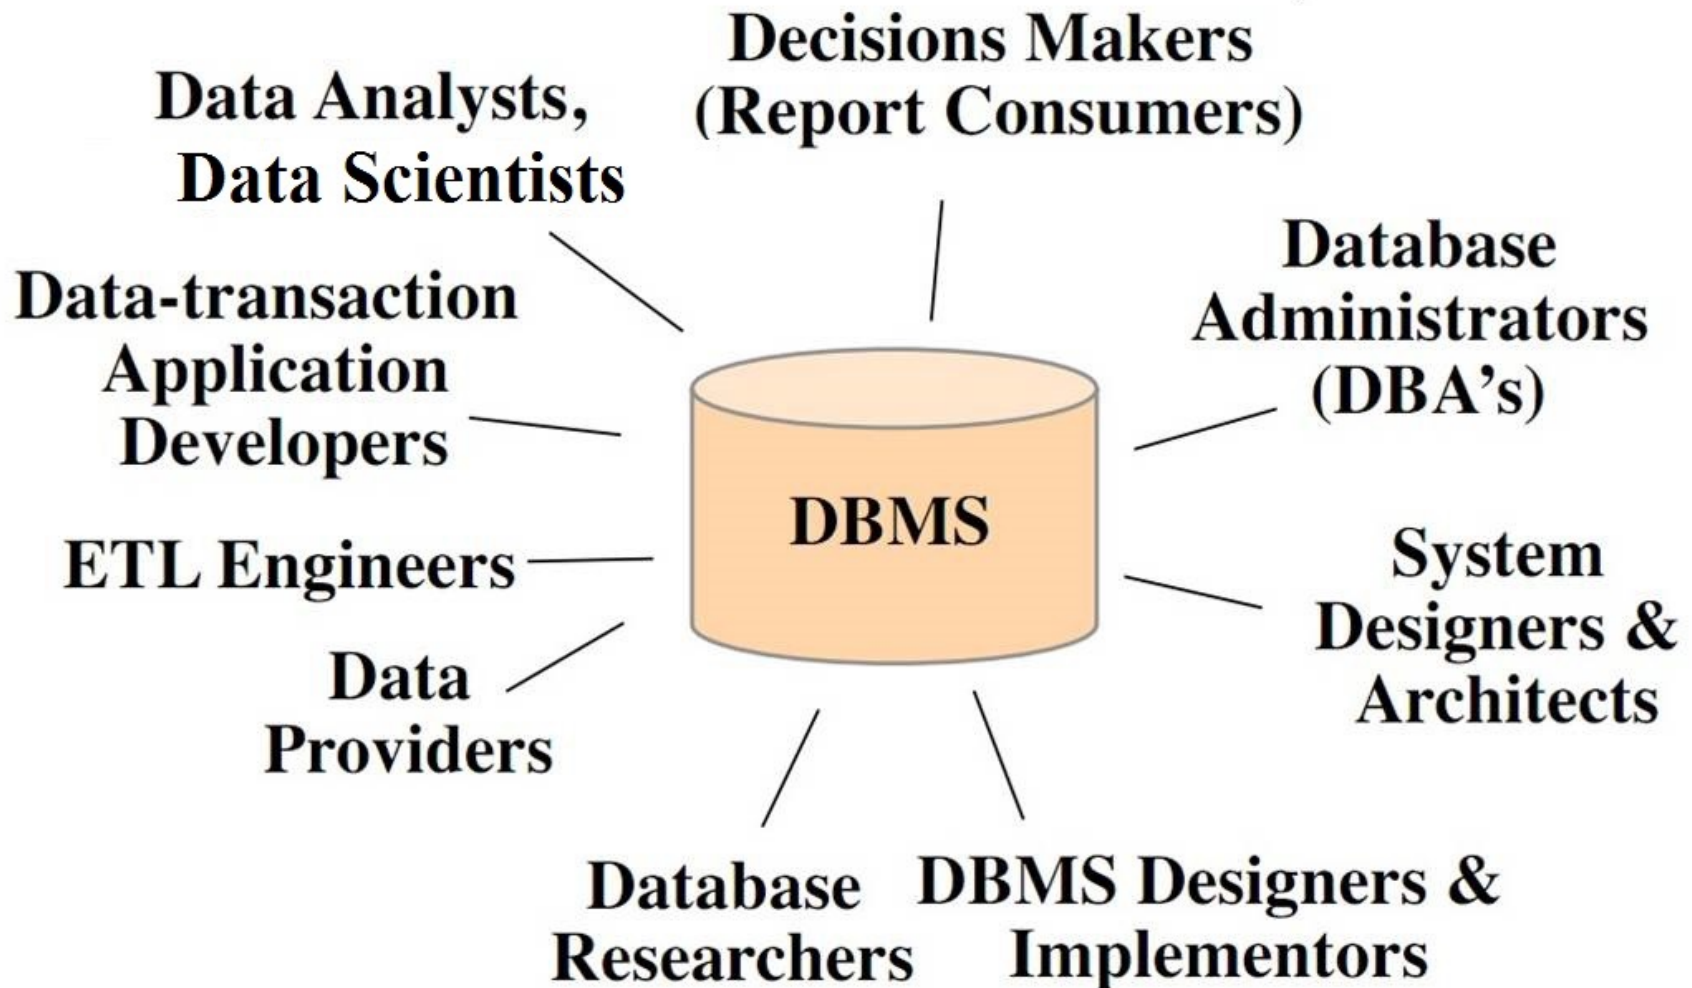
# History of Database Systems (continued)

- 2000s
  - Big data storage systems
    - Google BigTable, Yahoo PNUTS, Amazon Cloud
    - “**NoSQL**” systems: key-values, graph DBMSs, etc.
  - Big data analysis: beyond SQL
    - Map-reduce and friends
- 2010s and beyond
  - SQL “**reloaded**”
    - SQL is a front end to Map-Reduce systems (e.g. Hive)
    - Massively parallel database systems
    - Logical Data Warehouses
    - SQL-based multi-core in-memory databases
    - Spark SQL
- 2020s
  - Multi-Cloud Data Warehouses
    - E.g. Snowflake
  - Distributed Transactional Databases (“NewSQL” databases)
    - E.g. YugabyteDB





# Database Users: Who Needs to Use DBMS?



**Everybody who works with data!**

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access
  - **DDL:** Data Definition Language
  - **DML:** for Data Manipulation Language
    - Both use SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users

# Data Models

A collection of tools for describing

- Data
- Data relationships
- Data semantics
- Data constraints

## Types of Data Models

- Relational model
- Entity-Relationship (E-R) data model (mainly for database design)
- Object-based data models (Object-oriented and Object-relational)
- Semi-structured data model (e.g. XML)
- Other older models:
  - Network model
  - Hierarchical model



# Relational DBMS (RDBMS)

- Data stored in "tables" with relations between them.

**Students**

SID	Name
1	Alice
2	Bob
3	Charlie
4	Debra

**Courses**

CID	Name
1	ANLY640
2	ANLY502
3	ANLY503
4	ANLY511

**Enrollments**

Key

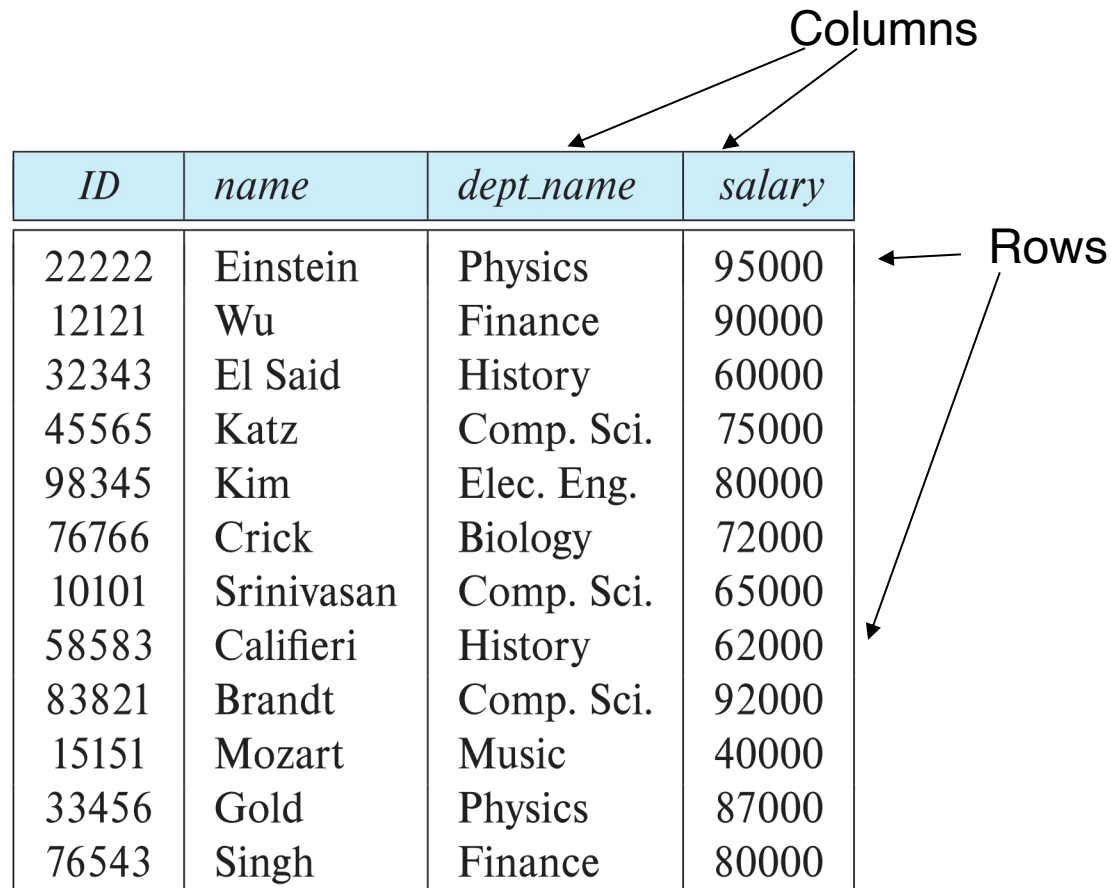
Foreign Keys

EID	SID	CID
1	1	1
2	1	2
3	2	1
4	3	2

- Alice is enrolled in ANLY640 and ANLY502
- Bob is enrolled in ANLY640
- Charlie is enrolled in ANLY502
- Debra is not enrolled in any course.

# Relational Model

- All the data is stored in various tables, called relations.
- Example of tabular data in the relational model



Columns

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

Rows

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users

# Database Design

The process of designing the general structure of the database usually involves the following steps

- Step 1: Requirements Analysis
- Step 2: Conceptual design
  - Database design requires that we find a “good” collection of relation schemas.
  - Goal: describe data and their relationships
- Step 3: Logical Design
  - Involves designing of the database schema.
  - Start of Implementation phase
- Step 4: Physical Design
  - Deciding on the physical layout of the database

# Requirements Challenge

- Data is collected much faster than it can be transformed into valuable information.
- Needed: identify the nuggets of information that is worth extracting from the avalanche of data – and what can be ignored.
- Catch 22:
  - IT department needs requirements to build a database
  - Customers often need to first understand the data in order to provide these requirements, which might not be possible until IT department builds an environment for them to use



*This is very important...  
Unimportant, of course, I  
meant—important,  
unimportant, unimportant,  
important.*

Lewis Carroll's Alice's  
Adventures in Wonderland

# Instances and Schemas

- Similar to types and variables in programming languages
- **Logical Schema** – the overall logical structure of the database
  - Example: University database consists of information about a set of students and courses in a university and the relationship between them
- **Physical schema** – the overall physical structure of the database, e.g. on disk
- **Instance** – the actual content of the database at a particular point in time (including data)

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- **Data Access (DDL and DML)**
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users

# Data-Definition Language (DDL)

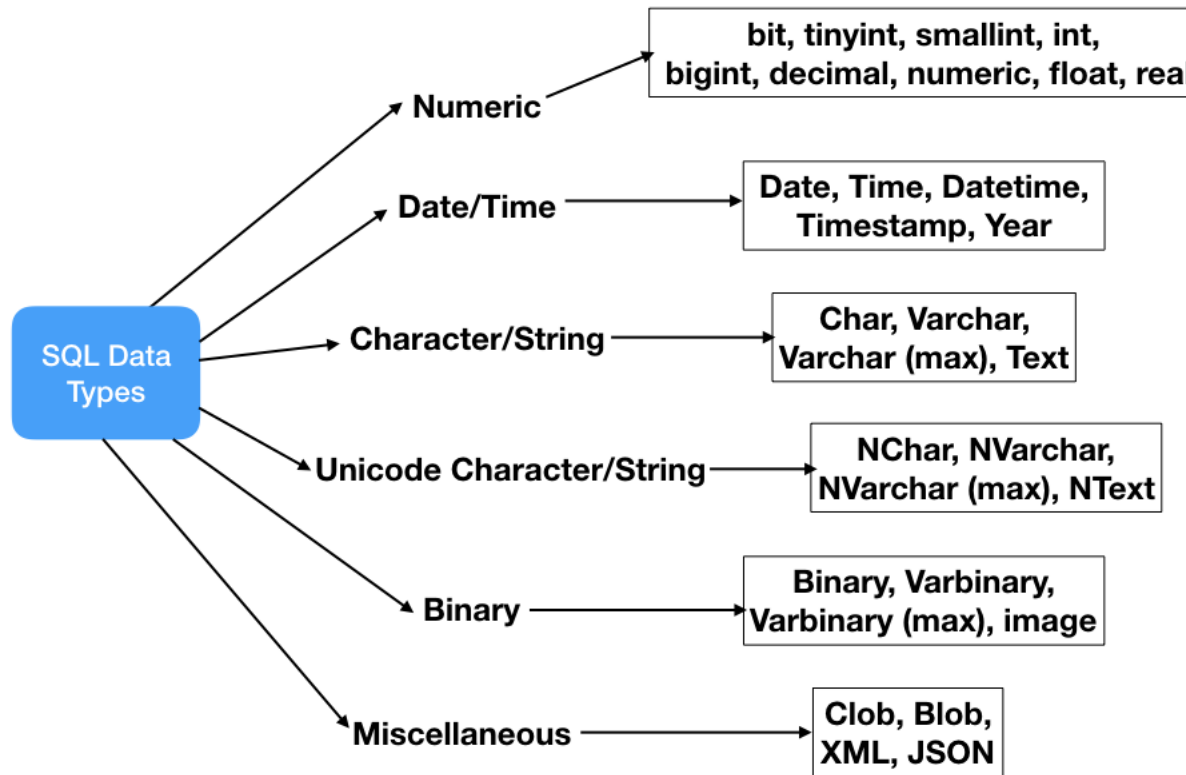
- Specification notation for defining the database schema

Example:   **create table** *instructor* (  
                              *ID*              **char**(5),  
                              *name*          **varchar**(20),  
                              *dept\_name* **varchar**(20),  
                              *salary*      **numeric**(8,2))

- DDL compiler generates a set of templates stored in a **Data Dictionary**
- Data dictionary usually contains:
  - Metadata (i.e., data about data)
  - Database schema
  - Constraints, e.g.
    - Primary key (e.g. ID uniquely identifies instructors) for referential integrity
    - User-defined (e.g. budget > 0)
  - Authorization: Who can access what



# Example of SQL Datatypes



<https://www.journaldev.com/16774/sql-data-types>

# Data Manipulation Language (DML)

- Language for accessing and updating the data
  - Allows to retrieve, insert, delete and update data organized by the appropriate data model
- There are two types of data-manipulation language
  - **Procedural DML** -- require a user to specify what data are needed and how to get those data.
  - **Declarative DML** -- require a user to specify what data is needed without specifying how to get that data.
- The portion of a DML that involves information retrieval is called a **query** language.
  - Though sometimes query language term is applied to all DML functions

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users

# SQL

- SQL is a standard language for querying and manipulating data
- SQL is a Declarative language
  - *What* but not *how*
- Many standards out there:
  - ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3), ....SQL:2016 (aka SQL8)
  - Vendors support various dialects and extensions

**SQL** stands for  
**Structured Query Language**

**Probably the world's most successful parallel programming language**

*Adapted from : <http://web.stanford.edu/class/cs145/>*

# SQL (continued)

- SQL commands — designed to read like English.
- Examples of SQL Commands are:
  - CREATE — Creates a database or table
  - SELECT — Reads data
  - INSERT — Inserts new data
  - UPDATE — Changes existing data
- Example: to find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```

# SQL Touches Every Piece of Modern Technology

The Oracle logo, featuring the word "ORACLE" in white capital letters on a red rectangular background.The Teradata logo, featuring the word "TERADATA" in orange capital letters.The Spark SQL logo, featuring the word "Spark" in black, an orange star icon, and the word "SQL" in black.

SQL on MarkLogic Server Quick Start

This chapter describes how to set up a typical development environment and a production environment on MarkLogic machines.

This chapter describes how to set up a typical development environment and a production environment on MarkLogic Server on separate machines.

The Apache Drill logo, featuring a stylized grey arrow pointing upwards and the words "APACHE DRILL" in blue capital letters.The Vertica logo, featuring the word "VERTICA" in blue capital letters with a stylized blue arrow pointing upwards.

Tenzing A SQL Implementation On The MapReduce Framework

The Research at Google logo, featuring the Google "G" logo and the text "Research at Google".The Tez logo, featuring the word "TEZ" in orange capital letters and a stylized orange flame icon.

SQL to MongoDB Mapping Chart

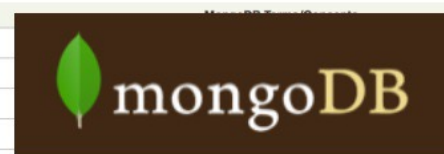
In addition to the charts that follow, you might want to consider the [Frequently Asked Questions](#) about MongoDB.

Terminology and Concepts

The following table presents the various SQL terminology and concepts and the corresponding MongoDB terminology and concepts.

SQL Terms/Concepts

database  
table  
row  
column  
index

The MongoDB logo, featuring a green leaf icon and the word "mongoDB" in a stylized font.

index

embedded documents and linking

primary key

In MongoDB, the primary key is automatically set to the `_id` field.

PostgreSQL



IBM DB2 - SQL Reference Vol 2 v8.2

The VoltDB logo, featuring the word "VoltDB" in blue and orange capital letters.The Greenplum logo, featuring a green leaf icon and the word "Greenplum" in black.

# Application Programs

- Non-procedural query languages such as SQL are not as powerful as a universal Turing machine.
- SQL does not support actions such as input from users, output to displays, or communication over the network.
- Such computations and actions must be written in a host language, such as C/C++, Java or Python, with embedded SQL queries that access the data in the database.
  - **Application programs** -- are programs that are used to interact with the database in this fashion.

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- **DBMS Engine**
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users



# DBMS Engine

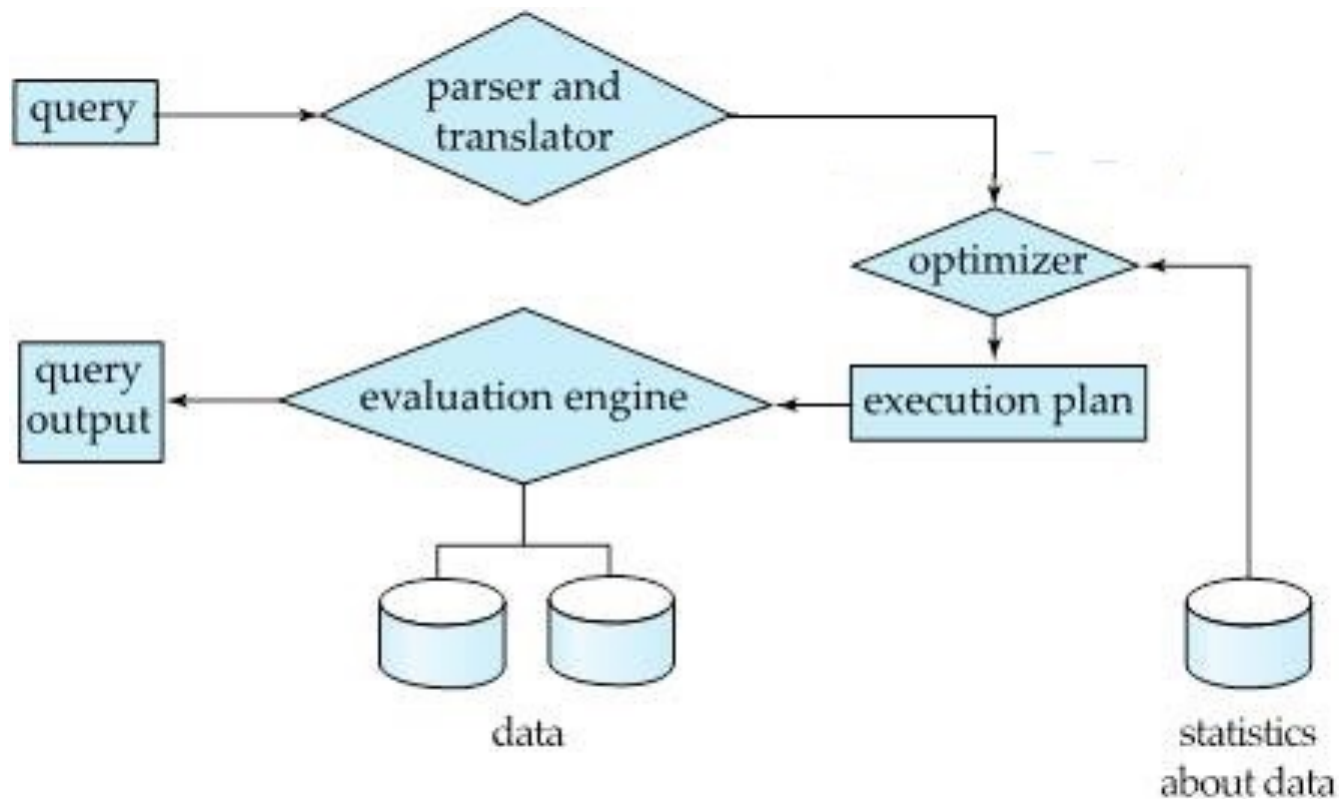
- A database system is partitioned into modules that deal with each of the responsibilities of the overall system.
- The functional components of a database system can be divided into:
  - **Query processor**
    - Helps to accelerate database processing
    - Can be done in parallel if more than one node/computer
  - **Transaction manager**
    - Allows users and developers to treat a sequence of db accesses as a single one
    - Manages concurrent access
  - **Storage manager**
    - Especially important for Massive Data storage
    - Can be done in parallel if more than one node/computer

# DBMS Query Processor

- The query processor components include:
  - DDL interpreter
    - Interprets DDL statements and records the definitions in the data dictionary.
  - DML compiler
    - Translates DML statements in a query language into an execution plan consisting of low-level instructions that the query evaluation engine understands.
    - The DML compiler performs query optimization; that is, it picks the lowest cost execution plan from among the various alternatives.
  - Query evaluation engine
    - Executes low-level instructions generated by the DML compiler.

# DML Query Processing

1. Parsing and translation
2. Optimization
3. Evaluation



# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - **Transaction manager**
  - Storage manager
- Database Architecture
- Database Users

# Transaction Manager

- **Transaction** is a collection of operations that performs a single logical function in a database application
- **Transaction-management component** ensures that the database remains in a consistent (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures.
- **Concurrency-control manager** controls the interaction among the concurrent transactions, to ensure the consistency of the database.

# Example of a Transaction

## Transaction Begins

```
UPDATE savings_accounts  
  SET balance = balance - 500  
  WHERE account = 3209;
```

Decrement Savings Account

```
UPDATE checking_accounts  
  SET balance = balance + 500  
  WHERE account = 3208;
```

Increment Checking Account

```
INSERT INTO journal VALUES  
  (journal_seq.NEXTVAL, '1B'  
   3209, 3208, 500);
```

Record in Transaction Journal

```
COMMIT WORK;
```

End Transaction

## Transaction Ends

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - **Storage manager**
- Database Architecture
- Database Users

# Storage Manager

- Storage manager is a program module that provides the interface between the low-level data stored in the database and the application programs and queries submitted to the system.
- The storage manager is responsible to the following tasks:
  - Interaction with the OS file manager
  - Efficient storing, retrieving and updating of data
- The storage manager components include:
  - Authorization and integrity manager
  - File manager
    - Manages allocation of space on disk
  - Buffer manager
    - Fetches data from disk into memory





# Storage Manager (continued)

- Raw data is stored on disk using the file system provided by OS
- The storage manager implements several data structures as part of the physical system implementation, for example:
  - Data files -- the database itself
  - Data dictionary -- metadata about the structure of the database, in particular, the database schema.
  - Indices -- can provide fast access to data items.
    - A database index provides pointers to those data items that hold a particular value.

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- **Database Architecture** (will start here next class)