

# L02: Database Design Using the E-R Model

DSAN 6300/PPOL-6810 : Databases Systems and SQL  
Programming

Irina Vayndiner

September 5 and 7, 2023



GEORGETOWN UNIVERSITY

# Logistics & Info for today's class

- Q01 will be published on **THU 9/7, due on TUE 9/19**
  - Watch Canvas for Announcements and due dates for assignments
- **Monday section: class on Thu 9/28 on zoom in lieu of 9/25 class**
- Use TA office hours (published on Canvas)
- Lecture
  - Database Architecture and Users
  - Database design using E-R model
- Lab: E-R model exercises

# Tentative Class Dates (updated)

	DSAN-6300-01 Mon	DSAN-6300-02& PPOL-6810 Thu
1	8/28	8/31
2	9/5 (Tue!)	9/7
3	9/11	9/14
4	9/18	9/21
5	9/28 (Thu, 9:30am) on zoom	9/28
6	10/2	10/5
7	10/16	10/12
8 (midterm)	10/23	10/19
9	10/30	10/26
10	11/6	11/2
11	11/13	11/9
12	11/20	11/16
13	11/27	11/30
14 (test)	12/5 (Tue, 10:30am)	12/5 (Tue, 10:30am)

## No classes

Mon, 9/4 Labor Day

Mon, 10/9: Mid Semester  
Holiday

Mon, 9/25: Yom Kippur -  
class on 9/28 instead!

Class added: Mon 12/5 for  
in lieu of 8/24 (all 3 sections)

# Tools Installation: AWS and MySQL Workbench

To install: **No Later Then (NLT) 9/22**

1. Install MySQL Workbench
2. Accept Invite to AWS Canvas
3. Attend TA office hours next week if you have questions!

**Special TA sessions on the week of 9/18 will be announced soon**

**!! You will not be able to use SQL in this Course without this setup!!**

# Outline of Today's Lecture

- Database Architecture and Users
- Database Design Using E-R Model
  - Overview of the Database Design Process
  - The Entity-Relationship Model
  - Mapping Cardinalities
  - Primary Key
  - Entity-Relationship Design Issues
  - Alternative Notations for Modeling Data
  - Other Aspects of Database Design

# Elements of DBMS Ecosystem

- Was discussed in Lecture 1
  - Data Models
  - Database Design
  - Data Access (DDL and DML)
  - SQL
  - DBMS Engine
    - Query Processor
    - Transaction manager
    - Storage manager
  
- Today:
  - Database Architecture
  - Database Users

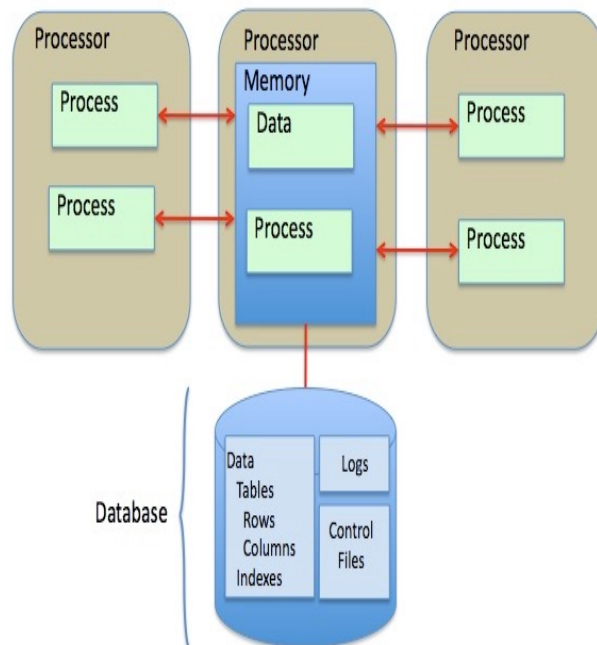
# Database Architecture Types

- Centralized (“shared”) databases
  - One to a few cores, shared-memory
- Client-server
  - One server machine executes work on behalf of multiple client machines.
- Multi-node databases can be
  - Shared-memory
  - Shared-disk
  - Shared-everything
  - Shared-nothing
- Distributed databases
  - Geographical distribution
  - Schema/data heterogeneity

# “Shared” RDBMS Architectures

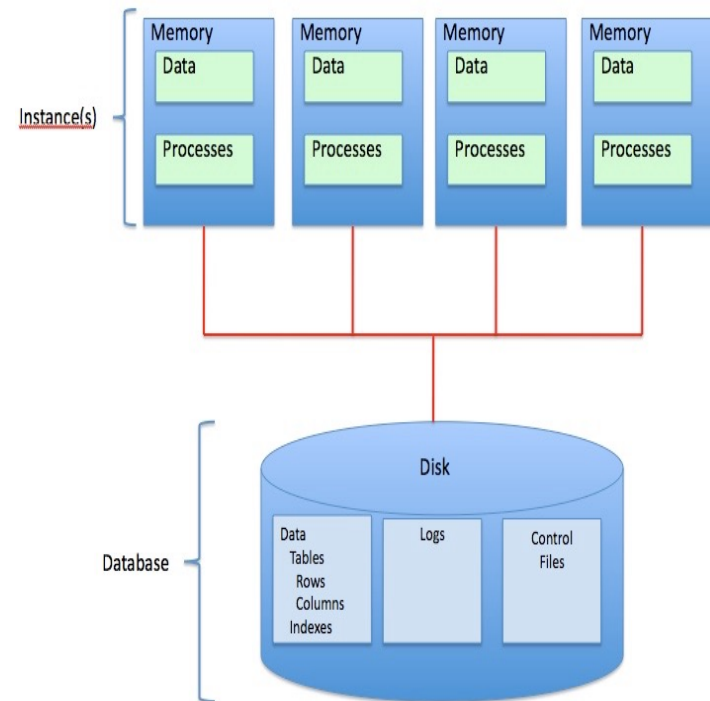
- SMP (symmetric multiprocessing). “Shared-everything” architectures share both disk and memory between nodes in the cluster. E.g.: IBM DB2 for z/OS, Sybase IQ.
  - Strength: Processing power
  - Limitation: Scalability

Shared-everything



“Shared-disk” E.g. Oracle RAC

Shared-disk



Ref: Wikipedia

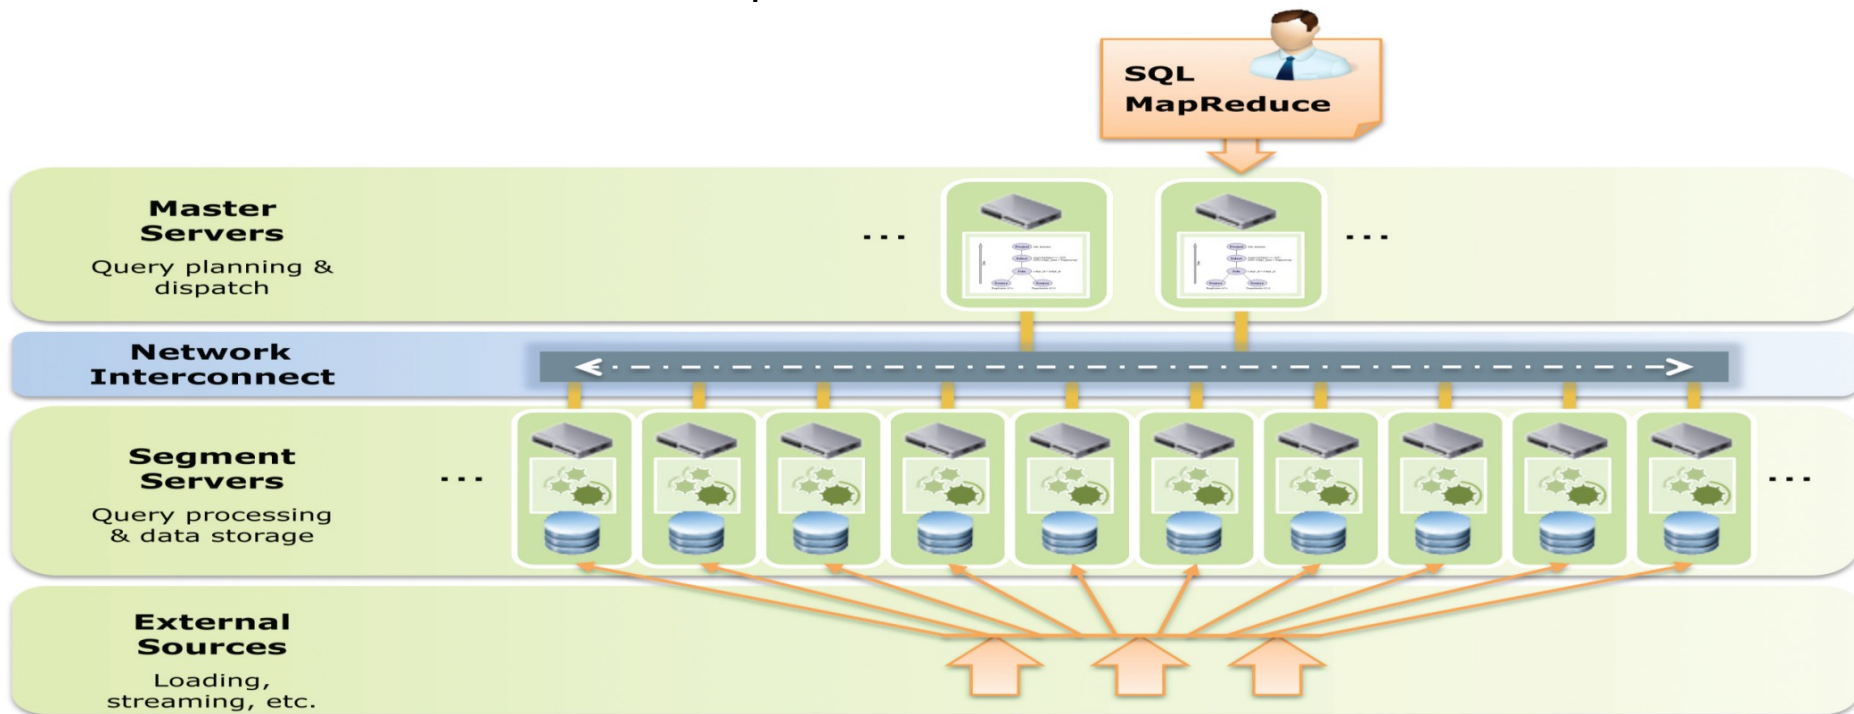


# “Shared-nothing” Architecture Massively Parallel Processing (MPP)

- Divide-and-conquer approach
- Massive refers to the concept of dozens or hundreds of nodes tied together achieving a single process
- Typical MPP architecture implements a shared-nothing (SN) paradigm where each parallel node operates self-sufficiently (“autonomously”) and controls its own memory and disk
  - A collection of nodes, each with local disk and local main memory, connected together by a high-speed network interconnect
  - Processor-disk pairs operating in parallel divide the workload to execute queries over large sets of data.
  - Adding nodes increases performance as well as capacity, scales linearly as new nodes are added into the system, to 100s of TBs and into PBs

# Shared Nothing (Massively Parallel) Relational Databases

- Nodes are connected by a high-speed Network Interconnect
- Master node acts as a “brain” that ties all the compute nodes together
  - Determines the optimal query plan
  - Dispatches that query to all the compute nodes
  - Gathers results from compute nodes and sends them to the user. All nodes



Source: GP Architecture White paper

# Elements of DBMS Ecosystem

- Data Models
- Database Design
- Data Access (DDL and DML)
  - SQL
- DBMS Engine
  - Query Processor
  - Transaction manager
  - Storage manager
- Database Architecture
- Database Users

# Database Users

There are four different types of database-system users

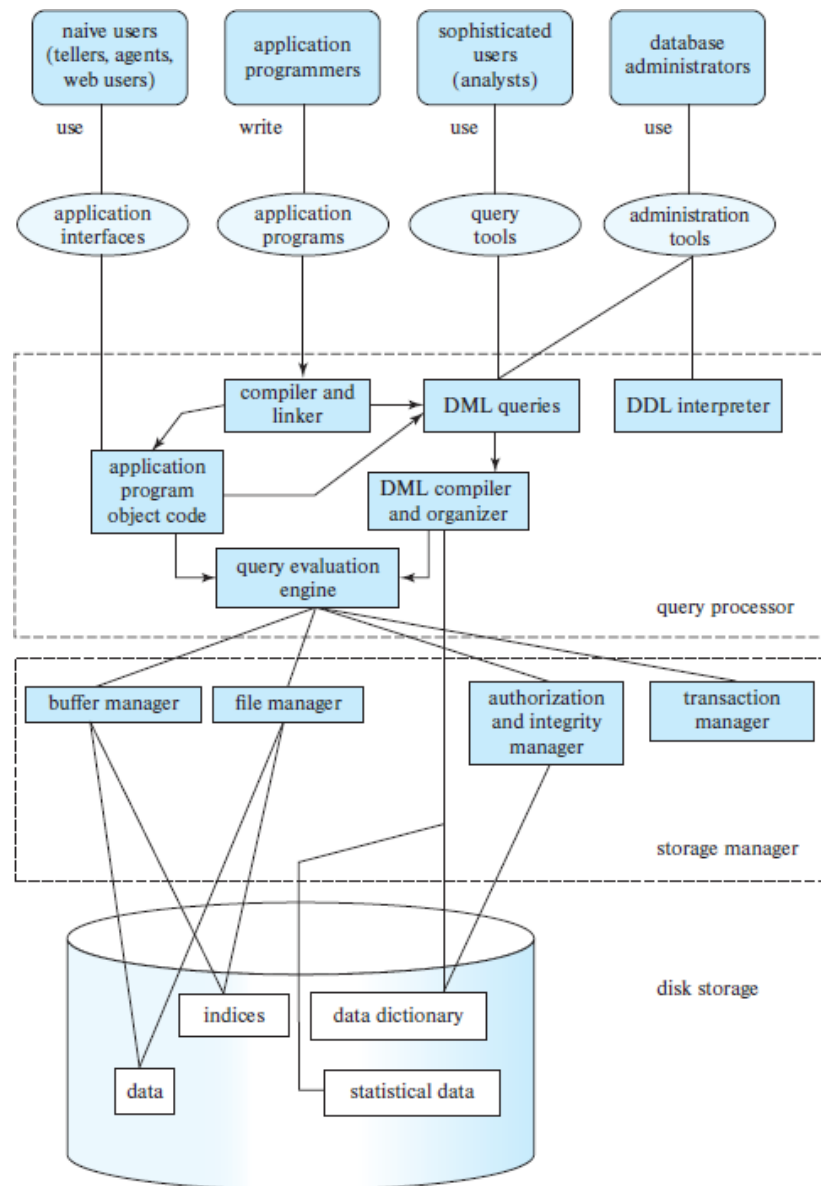
- **End users** -- users who interact with the system by invoking one of the application programs that have been written previously.
- **Application programmers** -- are computer professionals who write application programs.
- **Sophisticated users** -- interact with the system without writing programs
  - Using a database query language or
  - Using tools, such as data analysis software.
- **Specialized users** --write specialized database applications that do not fit into the traditional data-processing framework.
  - For example, CAD, graphic data, audio, video.

# Database Administrator (DBA)

A person who has central control over the DBMS is called a **database administrator (DBA)**, whose functions are:

- Schema definition
- Storage structure and access-method definition
- Schema and physical-organization modification
- Granting of authorization for data access
- Routine maintenance
- Periodically backing up the database
- Ensuring that enough free disk space is available for normal operations, and upgrading disk space as required
- Monitoring jobs running on the database and ensuring that performance is not degraded by very expensive tasks submitted by some users

# Putting It All Together



# Hands-on Exercises: MySQL and AWS RDS

- MySQL is the world's most popular open source relational database and Amazon RDS makes it easy to set up, operate, and scale MySQL deployments in the cloud.
- Amazon RDS for MySQL allows you up to focus on application development by managing time-consuming database administration tasks including backups, software patching, monitoring, scaling and replication.
- Amazon RDS supports MySQL Community Edition
  - We will use **8.0.34** in Class




# MySQL Ranking

## DB-Engines Ranking

The DB-Engines Ranking ranks database management systems according to their popularity. The ranking is updated monthly.

Read more about the [method](#) of calculating the scores.



<https://db-engines.com/en/ranking>

359 systems in ranking, August 2020

Rank			DBMS	Database Model	Score		
Aug 2020	Jul 2020	Aug 2019			Aug 2020	Jul 2020	Aug 2019
1.	1.	1.	Oracle +	Relational, Multi-model ⓘ	1355.16	+14.90	+15.68
2.	2.	2.	MySQL +	Relational, Multi-model ⓘ	1261.57	-6.93	+7.89
3.	3.	3.	Microsoft SQL Server +	Relational, Multi-model ⓘ	1075.87	+16.15	-17.30
4.	4.	4.	PostgreSQL +	Relational, Multi-model ⓘ	536.77	+9.76	+55.43
5.	5.	5.	MongoDB +	Document, Multi-model ⓘ	443.56	+0.08	+38.99
6.	6.	6.	IBM Db2 +	Relational, Multi-model ⓘ	162.45	-0.72	-10.50
7.	↑ 8.	↑ 8.	Redis +	Key-value, Multi-model ⓘ	152.87	+2.83	+8.79
8.	↓ 7.	↓ 7.	Elasticsearch +	Search engine, Multi-model ⓘ	152.32	+0.73	+3.23
9.	9.	↑ 11.	SQLite +	Relational	126.82	-0.64	+4.10
10.	↑ 11.	↓ 9.	Microsoft Access	Relational	119.86	+3.32	-15.47
11.	↓ 10.	↓ 10.	Cassandra +	Wide column	119.84	-1.25	-5.37
12.	12.	↑ 13.	MariaDB +	Relational, Multi-model ⓘ	90.92	-0.21	+5.96
13.	13.	↓ 12.	Splunk	Search engine	89.91	+1.64	+4.03
14.	↑ 15.	↑ 15.	Teradata +	Relational, Multi-model ⓘ	76.78	+0.81	+0.14
15.	↓ 14.	↓ 14.	Hive	Relational	75.29	-1.14	-6.51
16.	16.	↑ 18.	Amazon DynamoDB +	Multi-model ⓘ	64.75	+0.17	+8.18
17.	↑ 18.	↑ 25.	Microsoft Azure SQL Database	Relational, Multi-model ⓘ	56.85	+4.22	+28.85
18.	↓ 17.	↑ 20.	SAP Adaptive Server	Relational	53.96	+0.09	-1.90
19.	↑ 20.	↑ 21.	SAP HANA +	Relational, Multi-model ⓘ	53.12	+1.78	-2.31



# Some of MySQL Customers



# Amazon Relational Database Service (RDS)

<https://www.youtube.com/watch?v=eMzCI7S1P9M>

# Summary for Using Databases

- A major purpose of a Database system is to provide users with a way to manage and use that data.
- Underlying the structure of a database is the data model. The relational data model is the most widely deployed model for storing data in databases.
- A data-manipulation language (DML) is a language that enables users to access or manipulate data. SQL is the most widely used DML language.
- A database system has several subsystems: storage manager, query processor, transaction manager
- The architecture of a database system can be centralized or parallel

## Takeaways

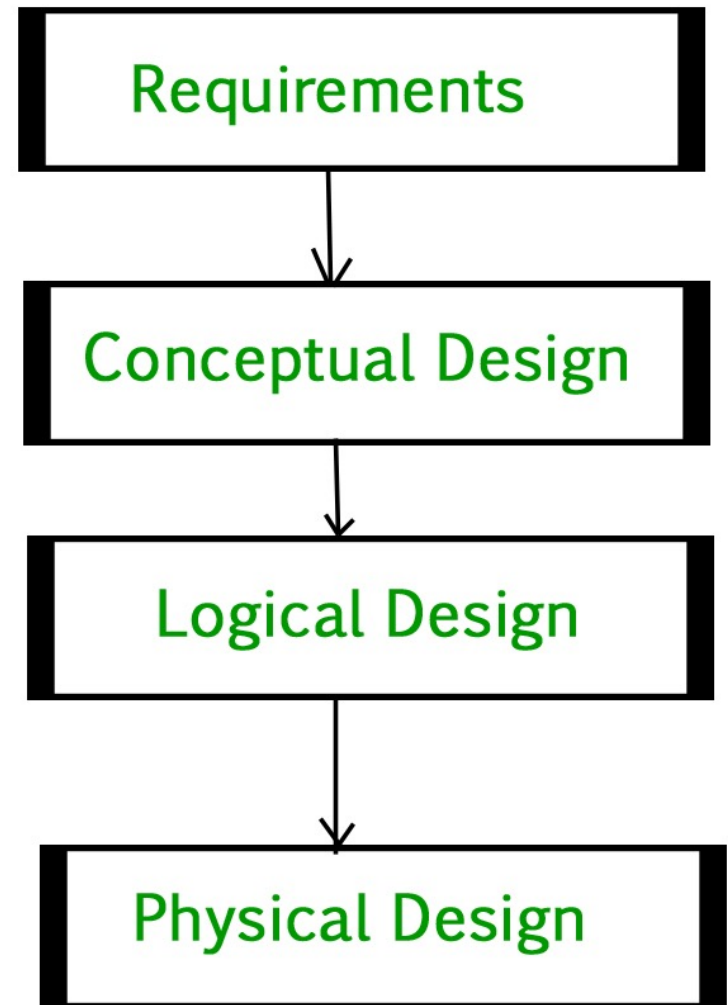
- Databases and SQL touch almost every part of modern technology
- This course will better equip you to succeed in your study, your research, and your job

# Today's main subject

## Database Design

# Database Design Phases

- Requirements Analysis-- characterize fully the data needs of the prospective database users.
- **Conceptual Design** -- choosing a **data model** (today's lecture!)
  - **The Entity-Relationship (E-R) model** is typically the result of conceptual design.
  - Conceptual schema specifies the **entities** that are represented in the database, the **attributes** of the entities, the **relationships** among the entities, and **constraints** on the entities and relationships
- Logical Design – Deciding on the database schema.
- Physical Design – Deciding on the physical layout of the database



# E-R model for Database Modeling

- The E-R data model was developed to facilitate database design by allowing specification of an **enterprise schema**
- The E-R data model employs three basic concepts
  - entity sets
  - relationship sets
  - attributes
- The E-R model also has an associated *graphical representation*, the **E-R diagram**, which can express the overall structure of a database graphically.

# Design Approach

- Entity-Relationship Model
  - Models an enterprise as a collection of *entities* and *relationships*
    - Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
      - Described by a set of *attributes*
    - Relationship: an association among several entities

# Definitions: Entity Sets

- An **entity** is an object that exists and is distinguishable from other objects.
  - Example: specific person, company, event, plant
- An **entity set** is a set of entities of the same type that share the same properties.
  - Example: set of all persons, companies, trees, holidays
- An entity is represented by a set of **attributes**; i.e., descriptive properties that are possessed by all members of an entity set.
  - Example:  
*instructor = (ID, name, salary)*  
*course = (course\_id, title, credits)*
- A subset of the attributes form a **primary key** of the entity set; i.e., **uniquely** identifying each member of the set.
  - We will discuss PK in detail later



# Entity Sets -- *instructor* and *student*

76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein



*instructor*

98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

*student*

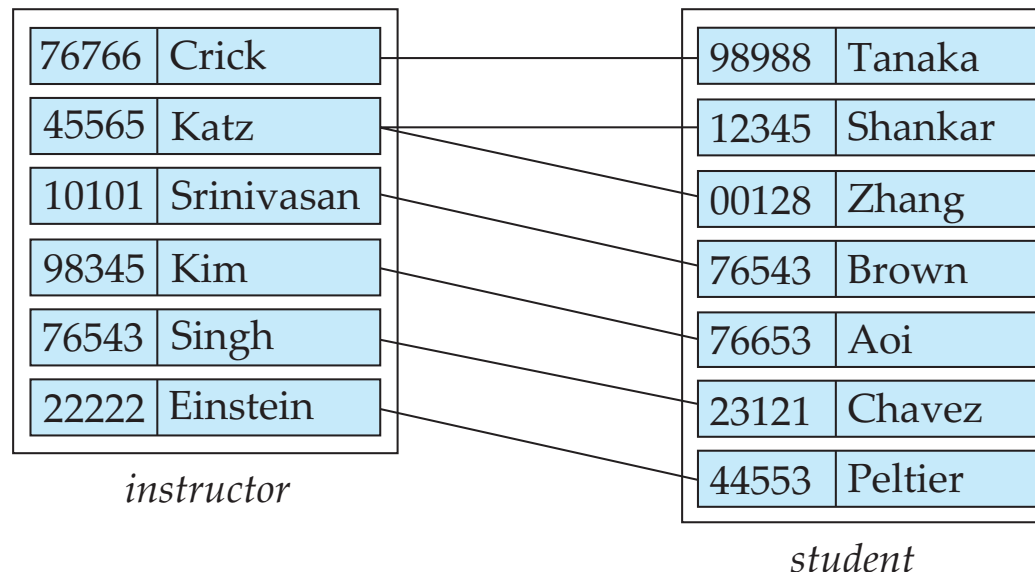
# Representing Entity Sets in E-R Diagram

- Entity sets can be represented graphically as follows:
  - Rectangles: represent **entity sets**
  - Attributes: listed inside entity rectangle
  - Underline: indicates **primary key** attributes




# Definitions: Relationship Sets

- A **relationship** is an association among several entities
  - Example: we define the relationship ***advisor*** to denote the associations between students and the instructors who act as their advisors.
- A **relationship set** is a collection of similar relationships – similar because it relates entities from the same entity sets
  - Pictorially, we draw a line between related entities.
  - Each line is a relationship, all lines are a relationship set



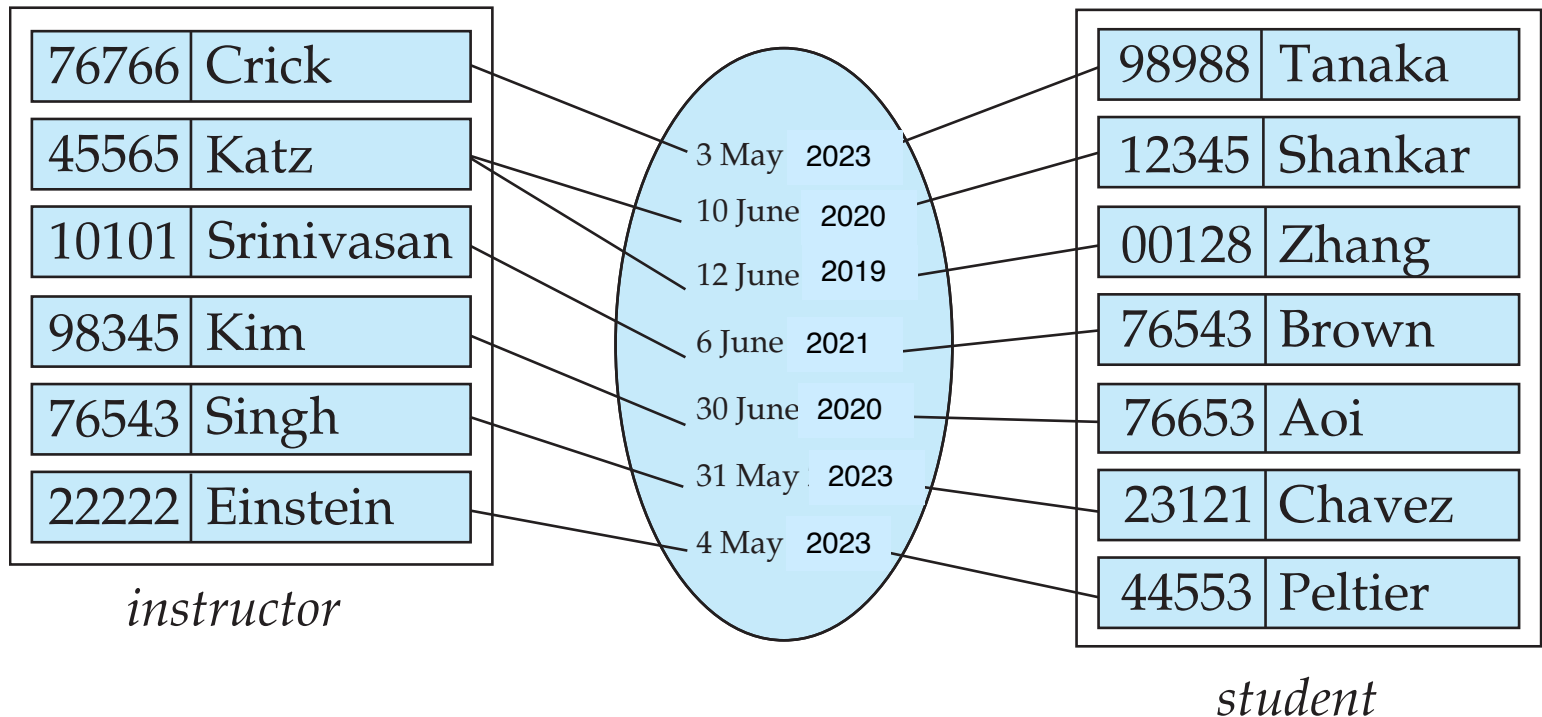
# Representing Relationship Sets via E-R Diagrams

Diamonds represent graphically relationship sets




# Attributes in Relationship Sets

- An attribute can also be associated with a relationship set.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the attribute *date* which tracks when the student started being associated with the advisor



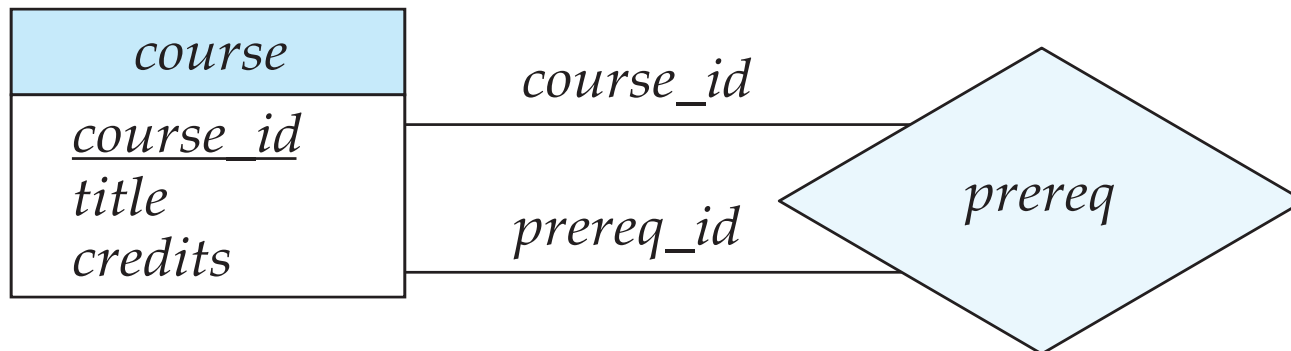
# Relationship Sets with Attributes

Dashed line represents the attributes of relationship sets




# Roles

- Entity sets of a relationship need not be distinct
  - Same entity set can participate in a relationship set more than once, in different roles
  - Each occurrence of an entity set plays a “role” in the relationship
- For example, let *course* be the entity set that records all courses in the University.
  - Say, Course 1 (C1) is a prerequisite for Course 2 (C2) => relationship set *prereq*
  - Ordered pairs of Course entities (C1, C2)
- The labels “*course\_id*” and “*prereq\_id*” are called **roles**.



# Degree of a Relationship Set

- Binary relationship
  - Involves two entity sets (degree = 2).
  - Most relationship sets in a database system are binary.
- Example of Relationships between more than two entity sets
  - Entity set *project* represents all research projects in the university
  - Each project can have multiple associated instructors and students
  - *students* work on research *projects* under the guidance of an *instructor*
  - Q: which instructor is guiding each student?
    - relationship *proj\_guide* is a **ternary** (degree =3) relationship between *instructor*, *student*, and *project*






# Attribute Types & Domain

- Attribute types
  - **Simple** and **composite** attributes, for example
    - Single attribute: name or address
    - Composite attribute: first\_name, mi, last\_name
  - **Single-valued** and **multi-valued** attributes, for example
    - Multivalued attribute: *phone\_numbers*
  - **Derived** attributes
    - Can be computed from other attributes
    - Example: adding *age*, given *date\_of\_birth* already exists
- **Domain (or value set)** – the set of permitted values for each attribute
  - Some examples:
    - *student\_name* can be a text string of a certain length
    - *semester* can have values of {Fall, Winter, Spring, Summer}


# Composite Attributes

Composite attributes allow us to divide attributes into subparts or sub-attributes (=component attributes).

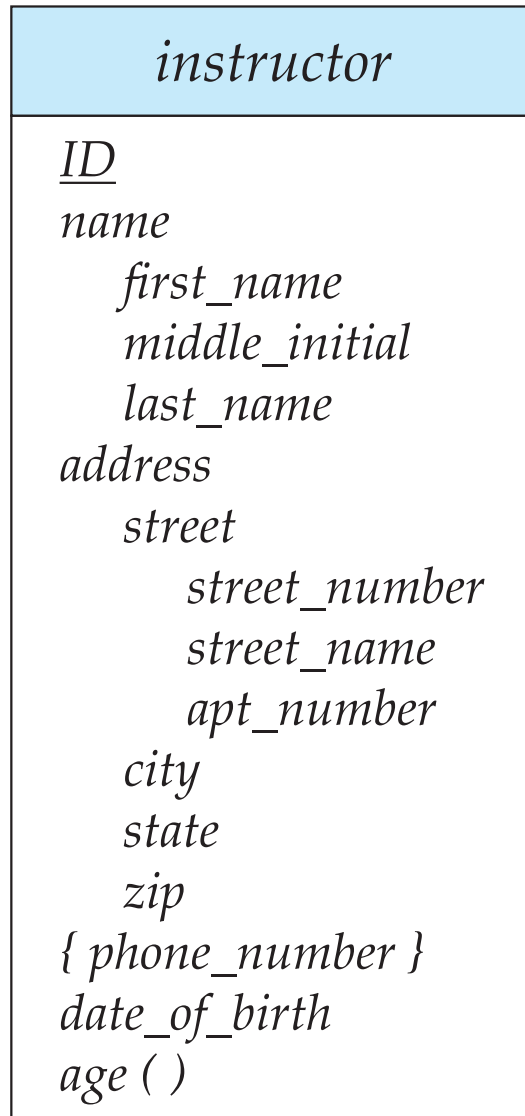
composite  
attributes



component  
attributes




# Representing Complex Attributes in E-R Diagram




# Mapping Cardinality Constraints

- **Mapping cardinalities** express the *number* of entities to which another entity can be associated via a relationship set. Most useful in describing binary relationship sets.
- For a binary relationship (between A and B) set the mapping cardinality must be one of the following types:
  - **One-to-one**
    - Entity in set A is associated with *at most* one entity in set B
  - **One-to-many**
    - Entity in set A is associated with *any number of* entities in set B
  - **Many-to-one**
    - Entity in set A is associated with *at most* one entity in set B
    - Entity in set B is associated with *any number of* entities in set A
  - **Many-to-many**
    - Entity in set A is associated with *any number of* entities in set B
    - Entity in set B is associated with *any number of* entities in set A

# Mapping Cardinalities




One to one




One to many

Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set

# Mapping Cardinalities



Many to one

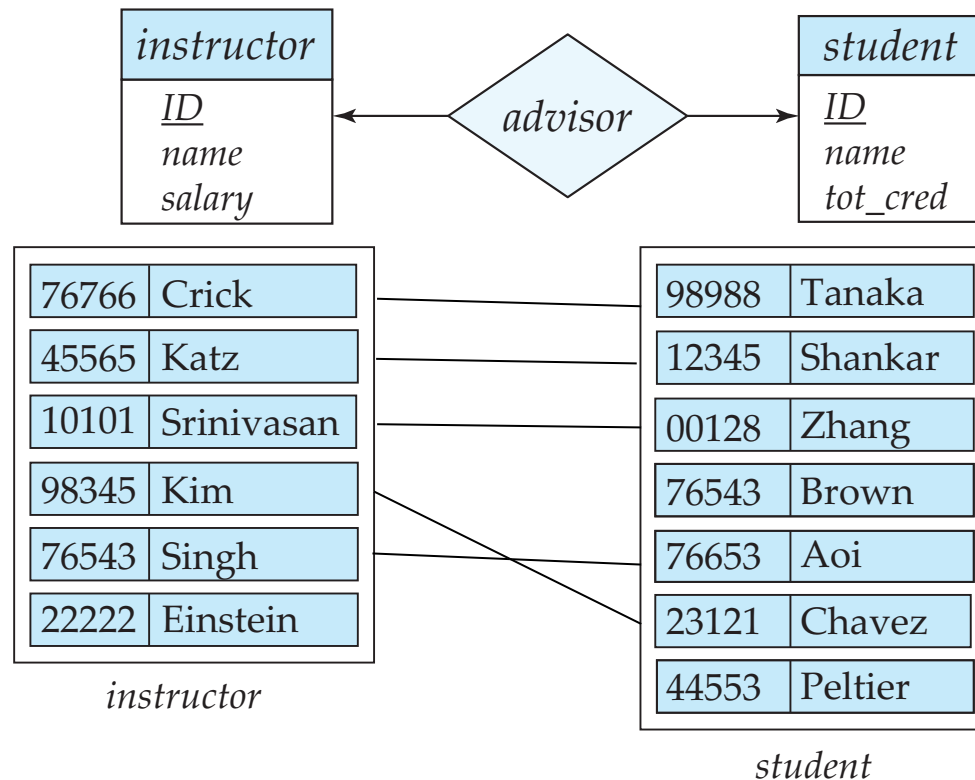


Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set

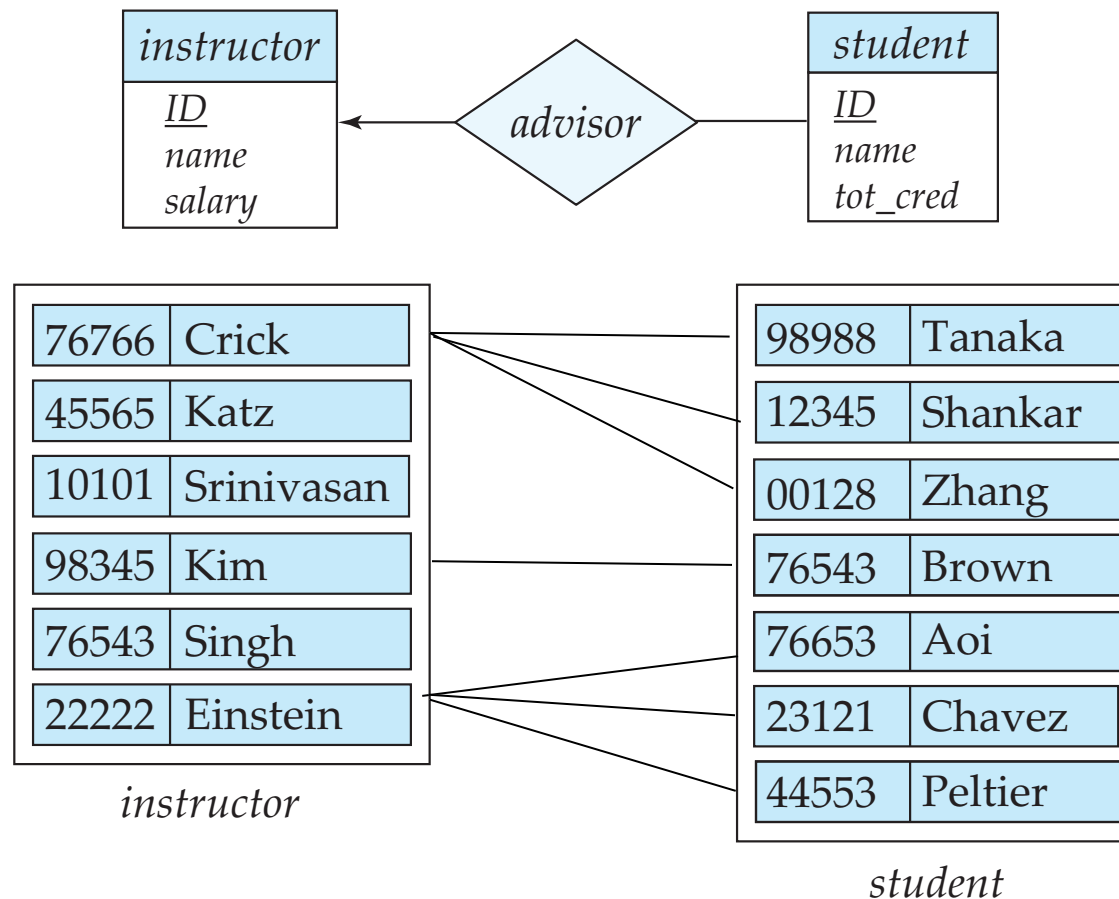
# Representing Cardinality Constraints in E-R Diagram

- We express cardinality constraints by drawing either a directed line ( $\rightarrow$ ), signifying “one,” or an undirected line ( $—$ ), signifying “many,” between the relationship set and the entity set.
- **One-to-one** relationship between an *instructor* and a *student*
  - A student is associated with at most one *instructor* via the relationship *advisor*



# One-to-Many Relationship

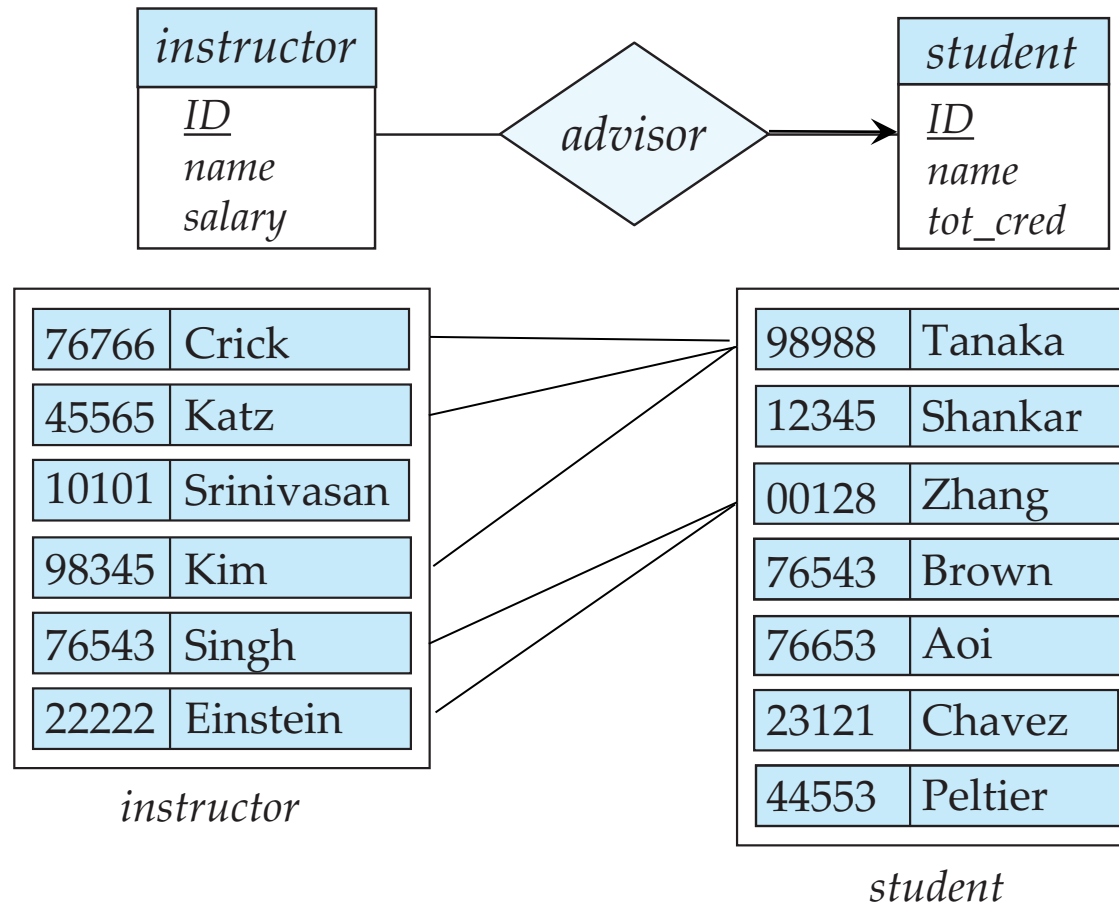
- **One-to-many** relationship between an *instructor* and a *student*
  - An instructor is associated with several (including 0) students via *advisor*
  - A student is associated with at most one instructor via *advisor*





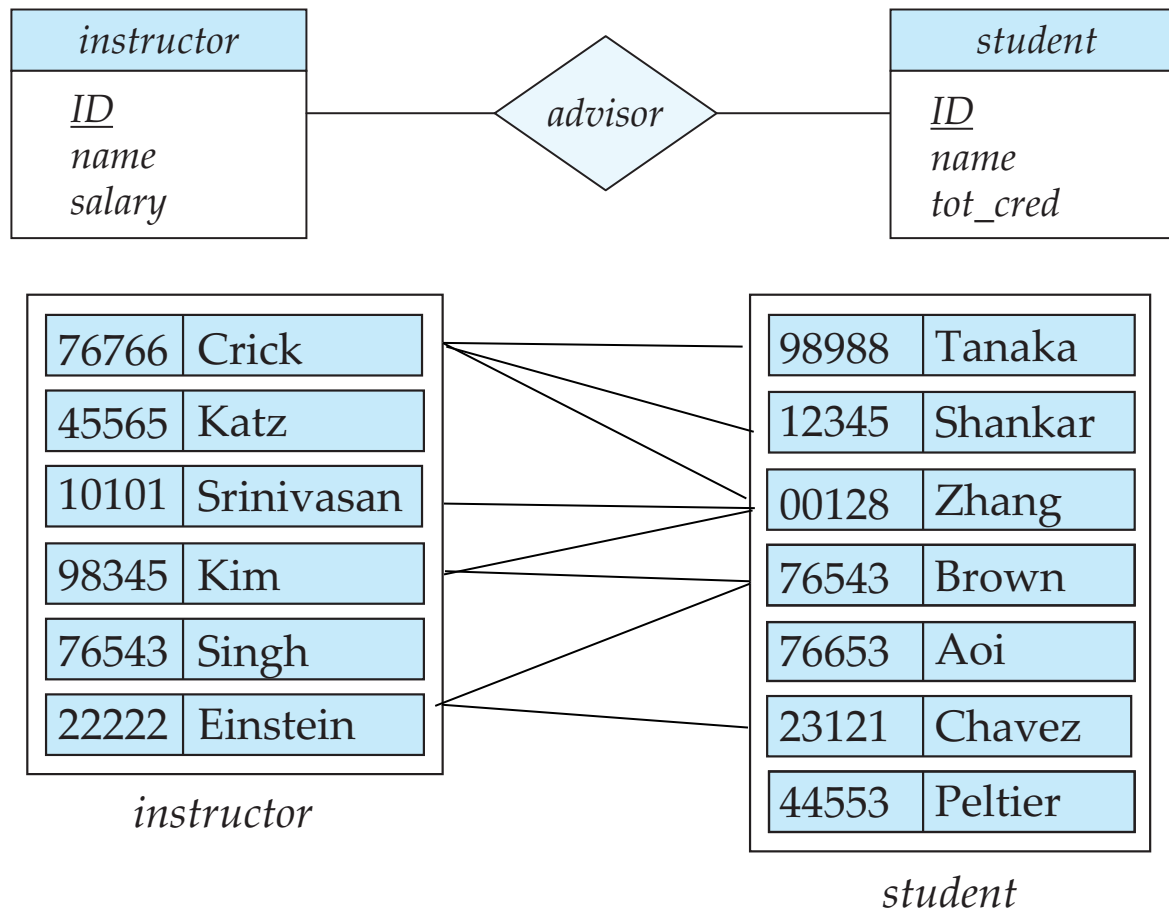
# Many-to-One Relationships

- In a **Many-to-One** relationship between an *instructor* and a *student*
  - An instructor is associated with *at most one* student via *advisor*
  - A student is associated with several (including 0) instructors via *advisor*



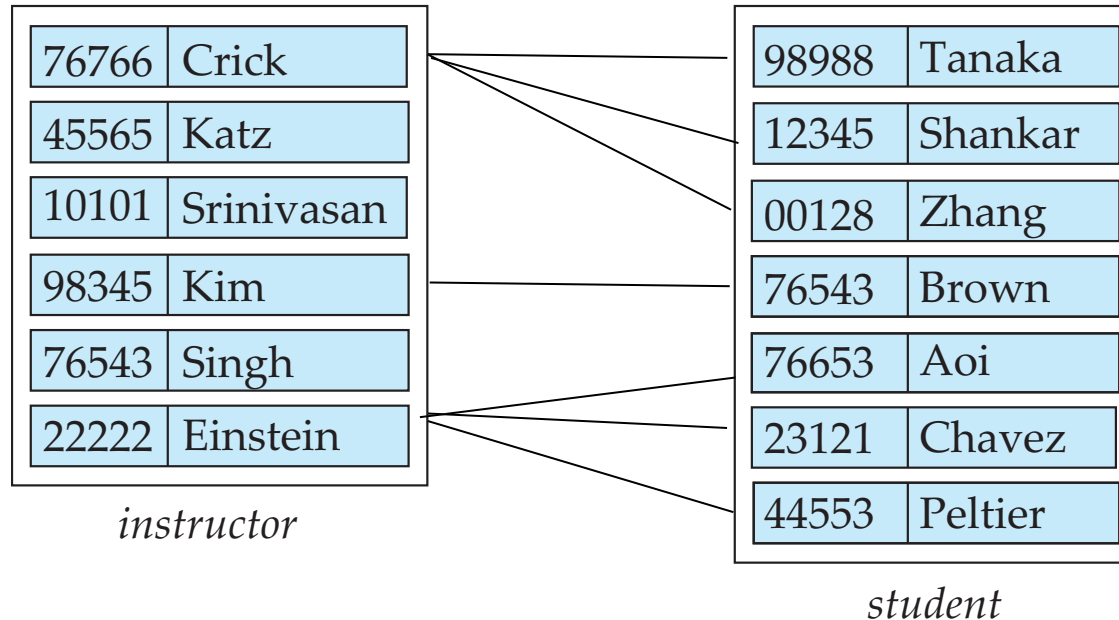
# Many-to-Many Relationship

- In a **Many-to-Many** relationship between an *instructor* and a *student*
  - An instructor is associated with several (possibly 0) students via *advisor*
  - A student is associated with several (possibly 0) instructors via *advisor*




# Total and Partial Participation

- **Total participation:** every entity in the entity set participates in at least one relationship in the relationship set
  - Example: participation of student in advisor relation is total
    - Every student must have an associated instructor
- **Partial participation:** some entities may not participate in any relationship in the relationship set
  - Example: participation of *instructor* in *advisor* is partial



# Total and Partial Participation (continued)


- **Total participation** indicated by *double line*.
  - Every entity in the entity set participates in at least one relationship in the relationship set
- **Partial participation** indicated by regular line.
- Some entities may not participate in any relationship in the relationship set



# Alternative Notation for Expressing Complex Constraints

- A line may have an associated minimum and maximum cardinality, shown in the form  $i.h$ , where  $i$  is the minimum and  $h$  the maximum cardinality
  - A minimum value of 1 indicates *total participation* (0 not allowed)
  - A maximum value of 1 indicates that the entity participates in at most one relationship
  - A maximum value of \* indicates no limit

- Example



- Instructor can advise 0 or more students. A student must have 1 advisor; cannot have multiple advisors



# Definitions: Key Constraints

- **Definition 1:** A set of attributes is a **candidate key** or a **key** of an entity set if it is both:
  1. **Unique:** No two entities can have same values for all key attributes
  2. **Minimal:** Cannot delete any of the attributes without impairing #1
    - Ex. for instructor( ID, name, dept\_name, email, ...)
      - {ID}
      - {email}
      - {name, dept\_name}
- **Definition 2:** A set of attributes is a **superkey** of an entity set if it is **unique** (no two entities can have same values for all key attributes), but the set is **not necessarily minimal**
  - Ex. instructor( ID, name, dept\_name, email, ...)
    - superkey: {ID, email}
- **THUS:** Any key is a superkey, but not the other way around

# Primary Key for Entity Sets

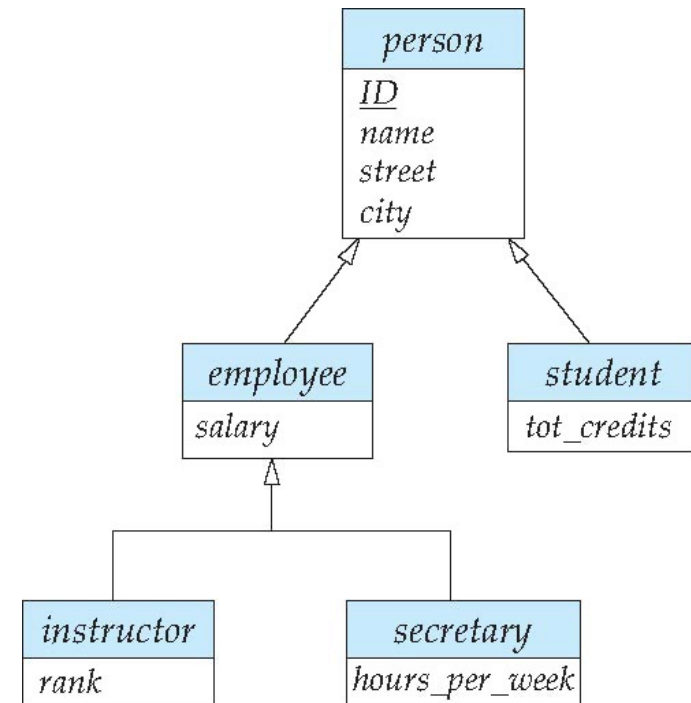
## ■ Primary Key (PK)

- One candidate key that you (or a DBA) pick to be the **primary key**
- It is important because -when translated into the relational model- the DBMS can enforce its uniqueness
  - You have to keep it unique over time
  - This takes work!
- The primary key is underlined in an E-R diagram (but not the candidate keys)



# Generalization

- **A bottom-up design process** – combine a number of entity sets that share the same features into a higher-level entity set.

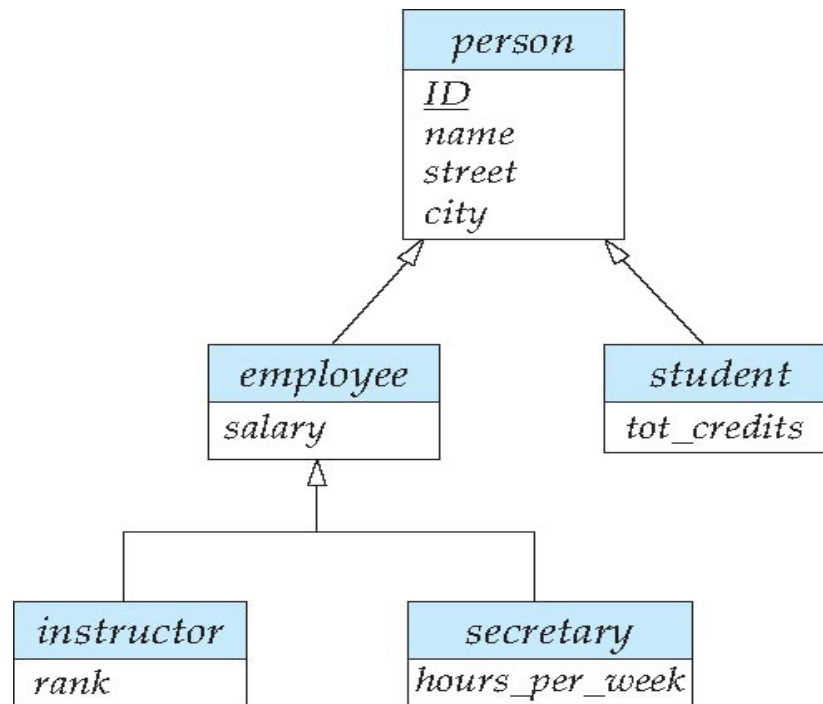




# Specialization

- Top-down design process; we designate sub-groupings within an entity set that are distinctive from other entities in the set.
- These sub-groupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set.
- Example, the entity set *person* may be further classified as one of the following:
  - *employee*
  - *student*

# Specialization Example




# Generalization & Specialization

- Specialization and generalization are simple inversions of each other
  - Represented in an E-R diagram in the same way. New levels of entity representation are distinguished (specialization) or synthesized (generalization)
  - The terms specialization and generalization are used interchangeably.
- Specialization and generalization may be total or partial type
  - **Total** - each high-level entity *must* belong to a low-level entity set
  - **Partial** - Some higher-level entities may not belong to any lower-level entity set.
- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked.


# Avoid Redundant Attributes

- Suppose we have entity sets
  - *instructor*
    - with attributes: *ID*, *name*, *dept\_name*, *salary*
  - *department*
    - with attributes: *dept\_name*, *building*, *budget*
- We model the fact that each instructor has an associated department using a relationship set *inst\_dept*
- The attribute *dept\_name* in *instructor* replicates information present in the relationship and is therefore *redundant* and needs to be removed.



# Aggregation

- One limitation of the E-R model is that it cannot express relationships among relationships.
  - Consider the ternary relationship *proj\_guide*, which we saw earlier
  - Suppose we want to record evaluations of a student by a project guide on a project




# Aggregation (continued)

- Issue: Relationship sets *eval\_for* and *proj\_guide* represent overlapping (redundant) information
  - Every *eval\_for* relationship corresponds to a *proj\_guide* relationship
  - However, some *proj\_guide* relationships may not correspond to any *eval\_for* relationships
    - So we can't discard the *proj\_guide* relationship
- Q: What can we do to get rid of that redundancy?
  - A: Eliminate redundancy via *aggregation*:
    - Treat relationship *proj\_guide* as an *abstract* entity
    - Allow relationships with that abstract entity
- Will show on the next slide

# Aggregation (continued)

- **Aggregation** as an abstraction through which relationships are treated as entities.
  - Thus, for our example, we regard the **relationship** set *proj\_guide* (relating the entity sets *instructor*, *student*, and *project*) as an **entity** set called *proj\_guide*.
- The following diagram represents aggregation:
  - A student is guided by a particular instructor on a particular project
  - A student, instructor, project **combination** may have an associated evaluation

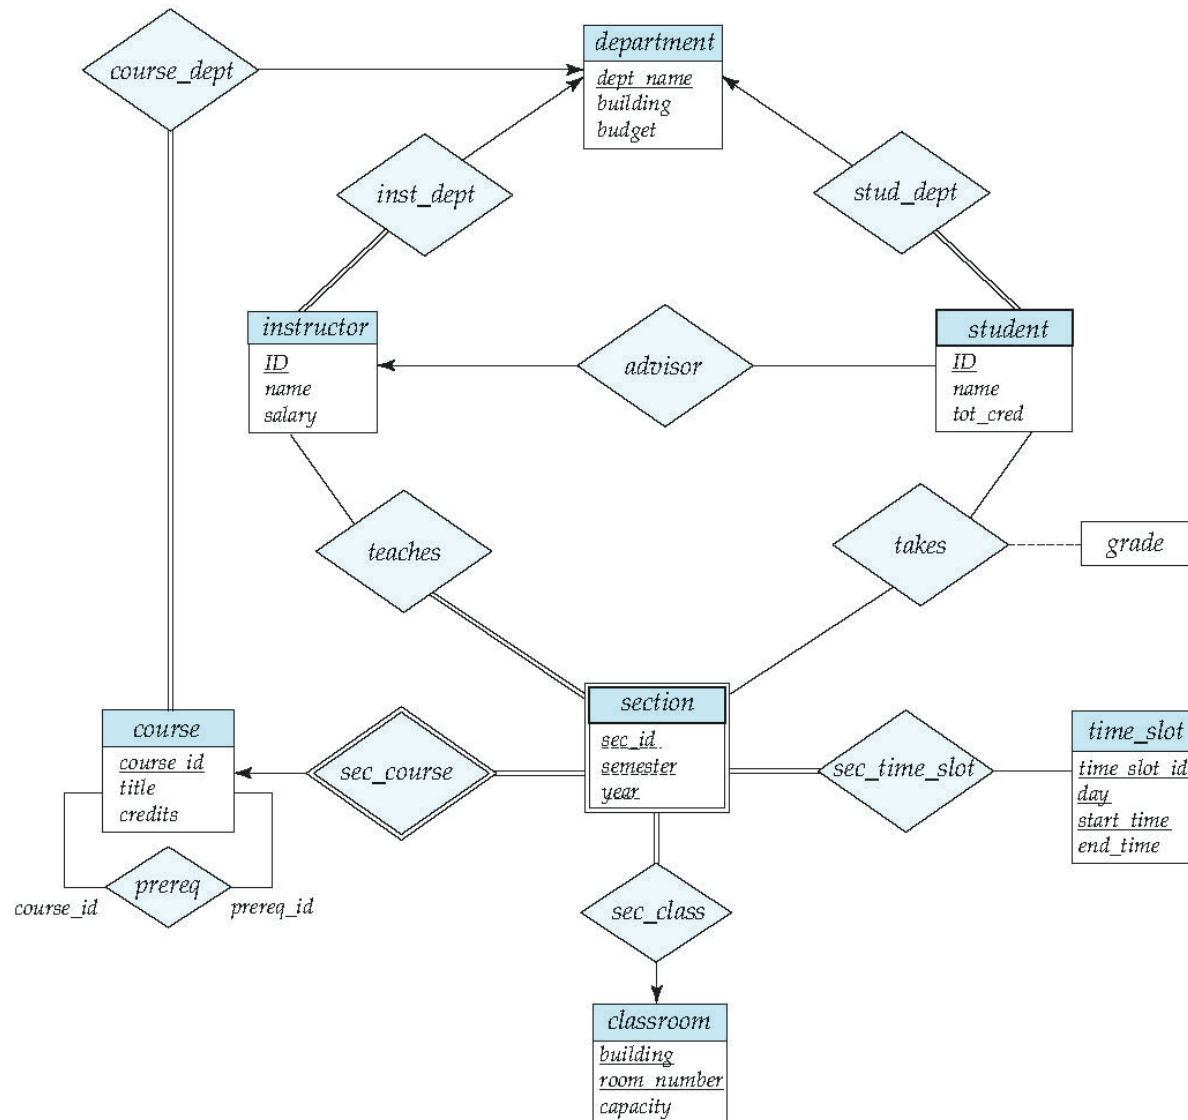


# University Database

- In this course we will be using a university database to illustrate all the concepts
  - It is published on Canvas
- Data consists of information about:
  - Students
  - Instructors
  - Classes
- Application program examples:
  - Add new students, instructors, and courses
  - Register students for courses, and generate class rosters
  - Assign grades to students, compute Grade Point Averages (GPA) and generate transcripts



# E-R Diagram for the University Database



# E-R Design Challenges

# Design Alternatives


- In designing a database schema, we must ensure that we avoid two major pitfalls:
  - Redundancy: a bad design may result in repeated information
    - Redundant representation of information may lead to data inconsistency among the various copies of information
  - Incompleteness: a bad design may make certain aspects of the enterprise difficult or impossible to model.
- Avoiding bad designs is not enough. There may be a large number of good designs from which we must choose.

# Design Choices

- Should a concept be modeled as an entity or an attribute?
- Should a concept be modeled as an entity or a relationship?
- Identifying relationships: Binary or ternary? Use Aggregation?
- Limits of the E-R Model:
  - A lot of data semantics & requirements can (and should) be captured
  - But some things just cannot be captured in E-R diagrams.

# Use Entities vs. Attributes


- Use of entity sets vs. attributes




- Should a phone number be modeled as an attribute or an independent (related) entity?
  - Phone number could be a character-string attribute of Instructor entities
  - Will not work in following cases
    - If we want to model that several instructors use the same phone, phone number must be a separate entity
    - If we want to support queries of phone number parts (e.g., area code, phone number can be modeled as a separate entity with sub-attributes)
    - If there can be multiple phone numbers per instructor, phone number must be an entity

# Erroneous Use of PK


- Common mistake 1: Use of Primary Key (dept\_name) of an entity set (department) as an attribute of a different entity set (student)



- Better way: dept\_name should not be an attribute of student, *relationship* stud\_dept is the correct way. It is always good to have relationship explicit.

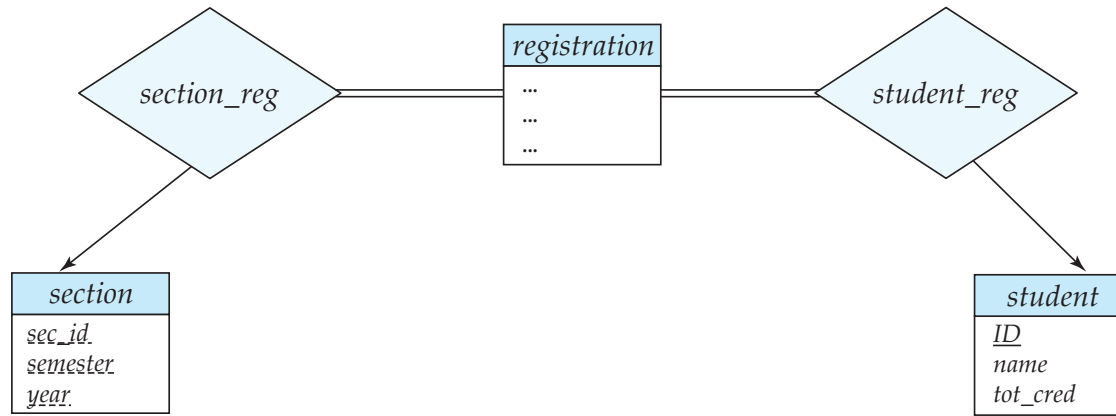


- Common mistake 2: To have both, attribute and relationship is redundant – should also be avoided



# Entities vs. Relationship sets

- Compare: *take* relationship in University diagram (shown earlier today) with introducing a new *registration* entity



- Considerations:
  - Pro: *take* relationship
    - More compact
    - Preferable in most cases
  - Pro: *registration* entity
    - Easier to add additional info to course-registration record
- Rule of thumb: designate a relationship set to describe an action (e.g. take) that occurs between entities

# Some Things E-R Just Doesn't Model Well

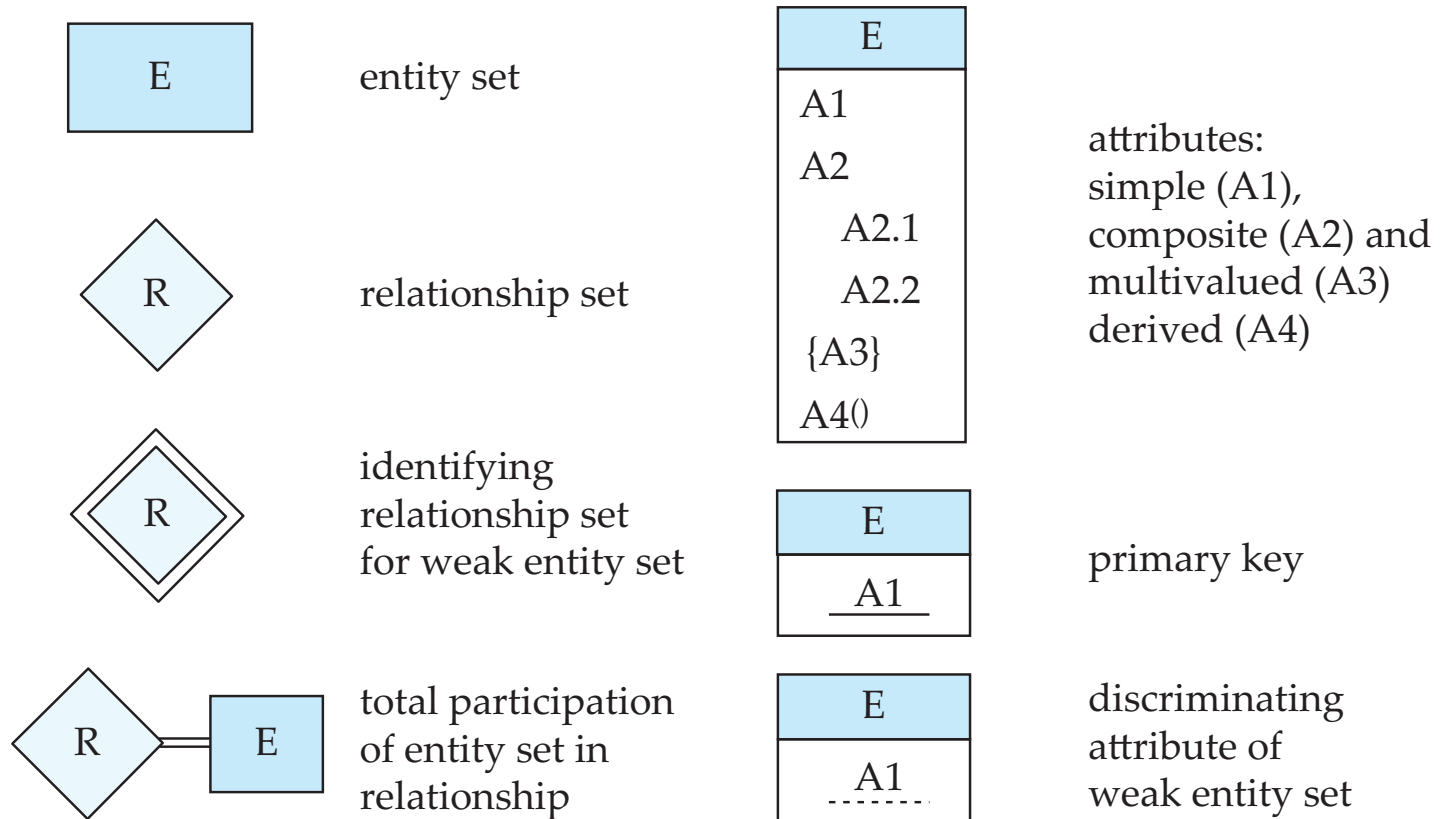
- Multi-valued attributes
- Numeric constraints
  - E.g. managers manage between 3 and 7 employees
- Functional dependencies
  - Fields that determine the value of other fields
    - e.g. that a department can't order two different parts from the same supplier
- Inclusion dependencies
  - Values in one attribute must be a subset of the values in another
- More general constraints
  - Managers must make 10% more than any of their employees
- All of that takes sometimes DBMS tools
  - Triggers/table constraints/application code!



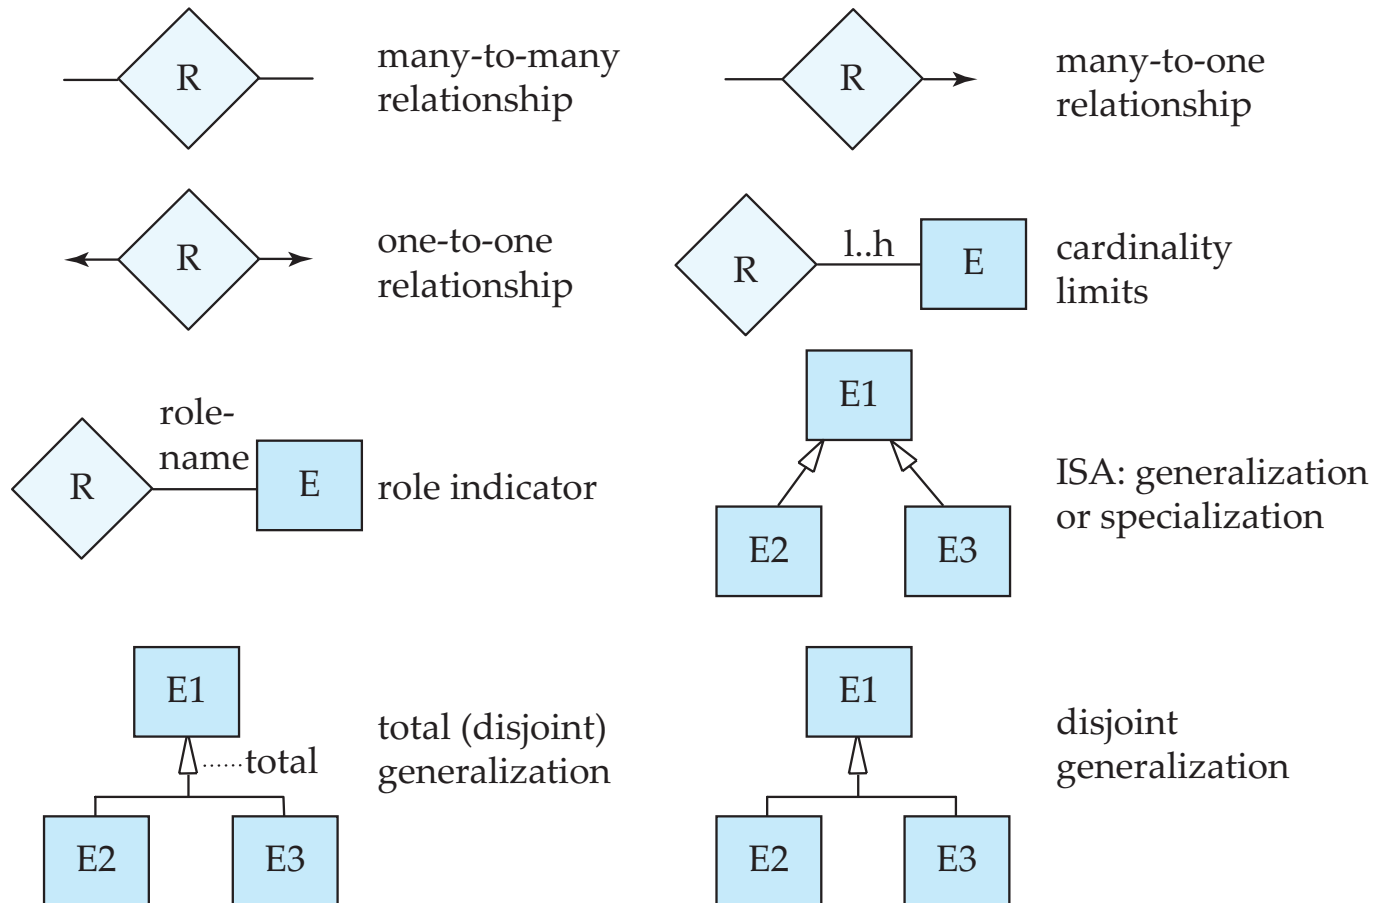
# Putting It All Together

- You attempt to capture all the described features in an E-R diagram
  - Determine entities, attributes, relationships
- You (almost certainly will) refine it several times, e.g.
  - moving attributes to the right place, making sure every constraint is correctly marked
  - ask clarifying questions as needed
  - test it out by thinking of example instances
  - explicitly state in English what you couldn't model in E-R
- The resulting E-R diagram has several important uses
  - you can translate it directly into a set of relations, which can be implemented during database creation stage (via SQL DDL statements)
    - (come back next week!)
  - ER diagrams are great documentation
    - succinct view of an entire information system
    - smart people (like you) will be able to pick one up and understand volumes!

# Summary of Symbols Used in E-R Notation



# Symbols Used in E-R Notation (continued)





# Alternative ER Notations



## Chen

## IDE1FX (Crows foot notation)



many-to-many  
relationship





one-to-one  
relationship



many-to-one  
relationship



participation  
in R: total (E2)  
and partial (E1)



# Lab: Developing E-R Diagram

**TO SUBMIT: part 2 of the Lab (only) upload to Canvas a Word Doc**

**Answer the following question:** What does this diagram enforce and what it does not enforce in relationships between package, sender and receiver.

**Grading rubric:**

- a. Full credit: explanation has to be understandable and make sense
- b. Partial credit: It does not make total sense
- c. Zero credits: not submitted or submitted a set of words that TA cannot understand after they read it 2 times