

L11: Data Warehousing, Data Analytics, and Transactions

DSAN 6300/PPOL 6810: Relational Databases and SQL Programming

Irina Vayndiner

November 9 and 13, 2023



GEORGETOWN UNIVERSITY

Classes left? Reminder: slide 14 from Lecture 1

	DSAN-6300-01 Mon	DSAN-6300-02& PPOL-6810 Thu
1	8/28	8/31
2	9/5 (Tue!)	9/7
3	9/11	9/14
4	9/18	9/21
5	9/25 (recording)	9/28
6	10/2	10/5
7	10/16	10/12
8 (midterm)	10/23	10/19
9	10/30	10/26
10	11/6	11/2
11	11/13	11/9
12	11/20 (on zoom)	11/16
13	11/27	11/30
14 (test)	12/5 (Tue, 10:30am)	12/5 (Tue, 10:30am)

No classes

Thu 11/24 Thanksgiving
Mon 12/4 – class on 12/5 instead

- HW4 (mini-project) is now published
 - 130 points
 - FAA data
 - Due **Wed, 12/6** (no extensions!)
 - If you submit by Mon 12/4, you will get 3 bonus points 😊
 - Plan accordingly
 - Multiple queries for one question on mini-project (only)? YES
 - Use Discussion Board
 - You will download your data set today
- After that one more assignment left:
 - Q04 to be available after the next class

Agenda for today's class

- Today:
 - Lecture: Data Warehousing, Data Analytics and Transactions
 - Lab: Rollups, etc.

Outline

- Data Warehousing and Data Analytics
- Online Analytical Processing (OLAP)
- Further information on Transactions

Overview of Data Analytics in RDBMS

- Using RDBMS for Data analytics
 - Used to make business decisions, e.g.
 - Per individual customer => Online Transaction Processing (OLTP)
 - E.g. what product to suggest for purchase to this customer
 - Across all customers, aggregated => OLAP (Online Analytical Processing)
 - E.g. what products to manufacture/stock, in what quantity
 - Both are important for businesses

Common Steps in Data Analytics

- Gather data from multiple sources into one location
 - Data warehouses sometimes integrate data into a common schema
 - Data often needs to be **extracted** from source formats, **transformed** to common schema, and **loaded** into the data warehouse
 - Can be done as **ETL (extract-transform-load)**, or **ELT (extract-load-transform)**
 - Example: AWS Glue <https://aws.amazon.com/glue/>



Common Steps in Data Analytics (continued)

- Generate aggregates and reports summarizing data
 - Dashboards showing graphical charts/reports
 - **Online analytical processing (OLAP) systems** allow dice & slice querying
 - Statistical analysis using tools such as R/Python/SAS/SPSS
 - Including extensions for parallel processing of big data

Common steps in data analytics (Continued)

- **Business Intelligence (BI)** is a “standard” Data Analytics
 - **Business intelligence** is "a set of methodologies, processes, architectures, and technologies that transform raw data into meaningful and useful information used to enable more effective strategic, tactical, and operational insights and decision-making.” Forrester Research
 - Often – descriptive (e.g. who did what when) or diagnostic (why is it happening)
 - The term **Decision Support (DS)** focuses on reporting and aggregation
 - Often includes Data Visualization
- **Machine learning (ML) analytics**



DATA WAREHOUSING

Data Warehousing

- Data sources often store only current data, not historical data
- Corporate decision making often requires a unified view of all organizational data, including historical data
- A **Data Warehouse(DW)** is a repository of information gathered from multiple sources, often stored under a unified schema of **RDBMS**
 - Examples:
 - AWS Cloud Redshift
 - Teradata
 - Oracle
 - MS SQL Server (also in the Azure Cloud)
 - Snowflake
 - Google Cloud BigQuery

DW Example: AWS Redshift

- Started as ParAccel
- Initial Redshift Release 2012
- RDBMS
- Based on Postgres
- Works well with many BI products, including data integration and visualization
- For Big Data Processing
- “Start small and scale up to Terabytes or Petabytes”


Example DW: Snowflake

- “Accelerate your analytics with the data platform built to enable the modern cloud data warehouse” Snowflake
 - Founded 2012, “DW as a Service”
 - The company's name was chosen as a tribute to the founders' love of winter sports
 - Cloud Storage
 - Automates data warehouse administration and maintenance
 - Runs on Amazon S3, on Microsoft Azure, and on the Google Cloud Platform
 - ANSI SQL compatible, with support for semi-structured data
 - Robust support for JSON-based functions
 - Optimized direct connectors for BI and Analytics tools
 - “Access charts and SQL analytics via Snowsight, the built-in visualization UI for Snowflake”.
 - Sept 2020, Snowflake made an IPO

Example DW: Google Cloud BigQuery

- Enterprise Data Warehouse in Google Cloud
- “BigQuery is a serverless, highly-scalable, and cost-effective cloud data warehouse with an in-memory BI Engine and AI Platform built in.” “You can focus on uncovering meaningful insights using familiar SQL without the need for a database administrator.” Google
- Has: Easy searchable query list, automated temp table expiration
- Built-in ML capabilities
- Google BigQuery queries need to be optimized to avoid high costs when pulling data
 - One of next lectures is on Database Optimization
 - Needs knowledge of SQL coding to leverage its data analysis capabilities.
 - You can do it! 😊
- More info: <https://cloud.google.com/bigquery>

Data Warehousing



Data Warehouse Design Issues

- *When and how to gather data*
 - **Source driven architecture:** data sources transmit new information to warehouse
 - Either continuously or periodically (e.g. at night)
 - **Destination driven architecture:** warehouse periodically requests new information from data sources
 - **Synchronous vs asynchronous replication**
 - Keeping warehouse exactly synchronized with data sources (e.g. using two-phase commit) is often too expensive
 - Usually OK to have slightly out-of-date data at warehouse
 - Data/updates are periodically downloaded from online transaction processing (OLTP) systems.
- *What schema to use*
 - Schema integration

Data Warehouse Design Issues (continued)

- **Data transformation** and **data cleansing**

- E.g. correcting inconsistencies in addresses (misspellings, zip code errors)
- E.g. **Merge** address lists from different sources and **purge** duplicates

- *How to propagate updates*

- Warehouse schema may be a (materialized) view of schema from data sources
 - View maintenance

- *What data to summarize*

- The raw data generated by a transaction-processing system may be too large to store online.
- However, we can answer many queries by maintaining just summary data obtained by aggregation on a relation
- For example, instead of storing data about every sale of clothing, we can store total sales of clothing by item name and category.

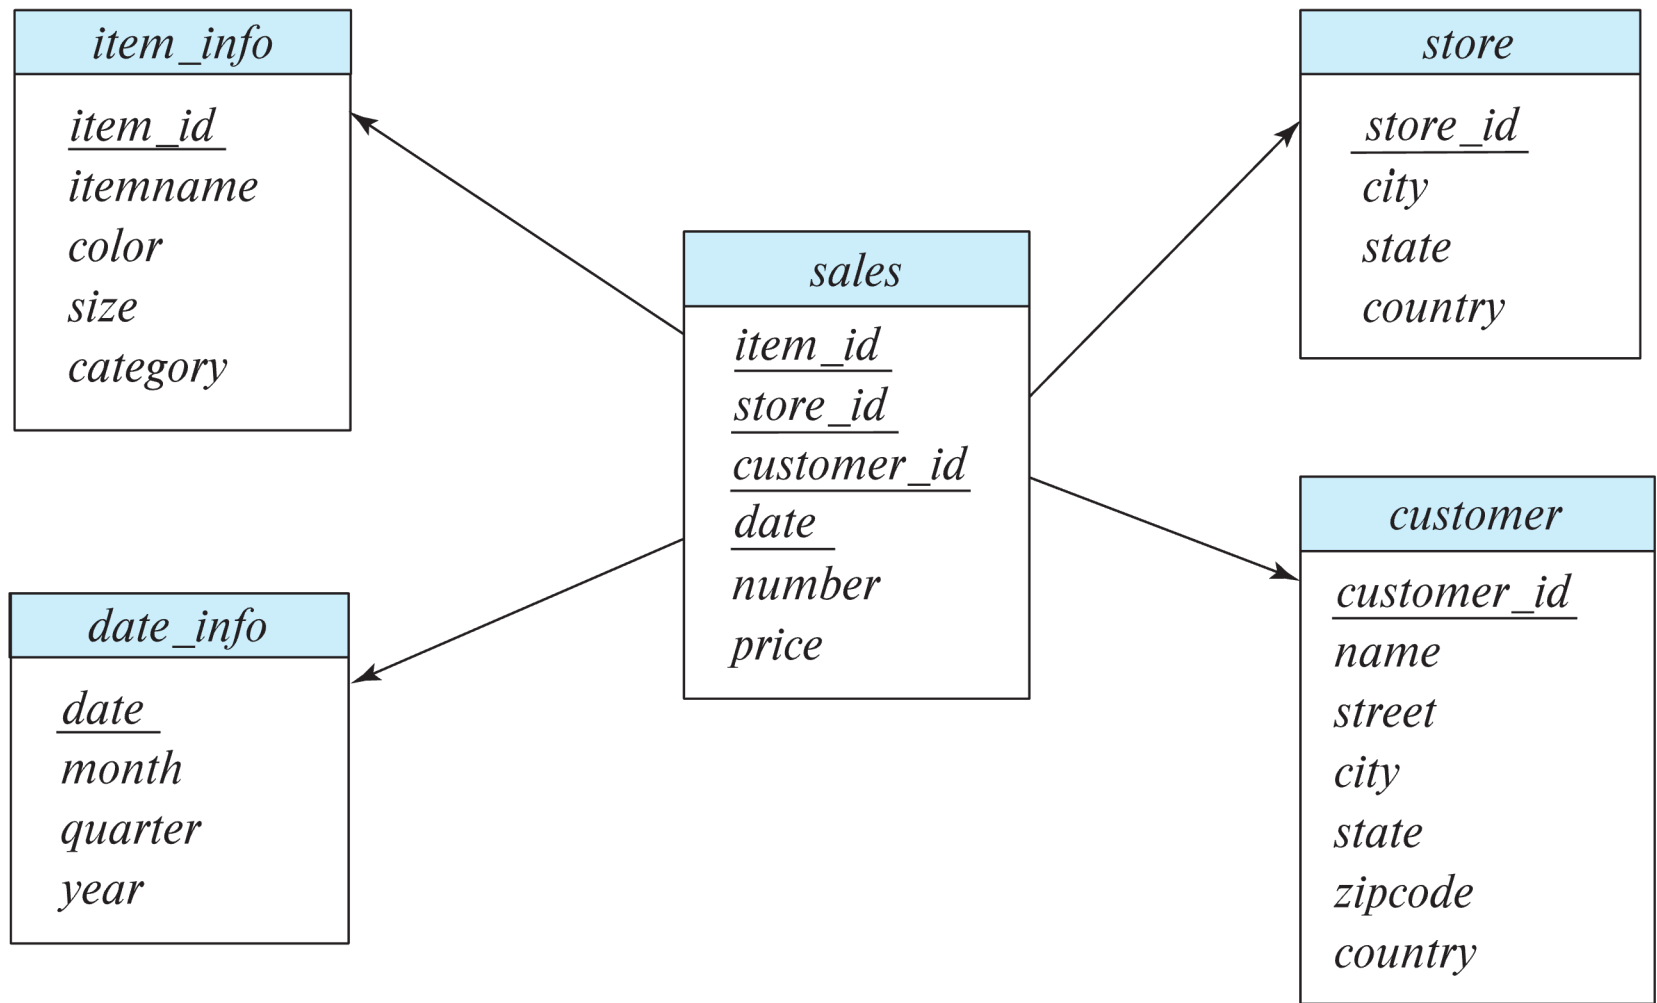
Multidimensional Data and Warehouse Schemas

- Important for HW4!
- Data in warehouses schemas can usually be divided into
 - **Fact tables**, which are large, e.g.
 - *sales(item_id, store_id, customer_id, date, number, price)*
 - **Dimension tables**, which are relatively small
 - Store extra information about stores, items, etc.
- Attributes of fact tables can be usually viewed as
 - **Measure attributes**
 - Measure some value, and can be aggregated upon
 - e.g., the attributes such as *number* or *price* of the *sales* relation
 - **Dimension attributes**
 - Dimensions on which measure attributes are viewed
 - e.g., attributes *item_id*, *color*, and *size* of the *sales* relation
 - Usually ids that are foreign keys to dimension tables


Star Schema

- Fact/dimension schema is called a **star schema**
 - More complicated schema structures
 - **Snowflake schema**: multiple levels of dimension tables
 - May have multiple fact tables
- Typically
 - Fact table joined with dimension tables
 - Need keys or star index for performance
 - Then aggregate on measure attributes of fact table

Data Warehouse Star Schema Example



- Some applications do not find it worthwhile to bring data to a common schema
 - **Data lakes** are repositories which allow data to be stored in multiple formats, without schema integration
 - Less upfront effort, but more effort during querying
 - A **Data swamp** is a deteriorated and unmanaged data lake that is either inaccessible to its intended users or is providing little value



- **Data Lakehouse:** newer data management paradigm
 - Combines the capabilities of data lakes and data warehouses
 - Enabling BI and ML on all data
 - Merging DW and Data Lake together into a single system means that data teams can move faster as they are able to use data without needing to access multiple systems.

Database Support for Data Warehouses

- Data in warehouses usually append only, not updated
 - Can avoid concurrency control overheads
- Data warehouses often use **column-oriented storage**
 - Columns are compressed, reducing storage, IO and memory costs significantly
 - Queries can fetch only attributes that they care about, reducing IO and memory cost
- Data warehouses often use **parallel** (MPP) storage and query processing infrastructure

1. Use Cloud Data Warehouse (vs On-Prem)

- Flexible and scalable environment
 - Minimal setup and maintenance (e.g. replication, backups, etc)
 - Storage and compute resources suitable for Big Data scale
 - On demand, automatic scalability is useful for highly unpredictable, ad-hoc query workloads that result from self-service practices.
 - Separates computing from storage
 - One can provision extra computing for the duration of batch job only
 - e.g. on weekends
 - Easy to integrate with other cloud services
 - Can be lower cost
- Cloud Models
 - Fully in the cloud
 - Hybrid (stable part “on prem”, dynamic in the cloud)
 - Multiple-clouds DW

2. Make Use of Data Variety

- Modern data warehouses support a wide range of data types and analytics
 - Structured is still a top data source, e. g. from multiple RDBMSs
 - Semi- (JSON, XML) and un-structured (text, multimedia) data from Data lakes
 - Files from mainframes
 - Geospatial data
 - Sensors and IoT devices
- Data generated in the cloud, e.g. demographics, weather, home sale prices, etc.
 - Analyze data where it lives, aka “data gravity”
 - Moving large volumes of data is resource intensive

3. Use ELT (vs ETL)

- Data is first extracted from multiple sources, loaded, then transformed
 - Transformation can use parallel processing power of the cloud platform
 - Data loading tools use SQL
 - Workflows can be saved and re-used

4. Integrate tools for data discovery and analytics

- Movement toward self-service
 - Many self-service tools can enable immersive, “speed-of-thought” experiences with data and analytics.
 - Modern self-service platforms often provide external integration with R and Python.
- Automation
 - E.g. Data preparation, finding Insights
 - Machine learning infused into the cloud data warehouse to help with optimization

5. Data Governance

- Set of processes, roles, standards, and measures that ensure important data assets are formally and consistently managed throughout the enterprise
 - Trusted, curated environment for Data
 - Data quality, availability, usability, consistence, integrity, and security
- For Data Governance in the Cloud the following is needed
 - Visibility into Cloud, including authentication attempts, queries ran, etc
 - Cloud **metadata** that helps with data consistency
 - Where the data was created, who owns it, what it is about, what it is used for, how it is organized, where it is located
 - Data Catalog, including tables, indices, views, etc.
 - Data Provenance or Lineage
 - » Where data originated and how it has been changed and transformed in Data Warehouse

6. Integrated Analytics Stack

- Requirements, functionality
 - E.g. how many data sources, what kind of analytics
- Cost considerations
 - Based per Time/Query/Cluster(?)
 - Tune your queries!
 - Plan for how much data you are pulling
 - Tight integration with other tools including security tools

OLAP

■ Online Analytical Processing (OLAP)

- Interactive analysis of data, allowing Big Data data to be summarized and viewed in different ways with negligible delay
 - Analyze multidimensional data from multiple perspectives
- Optimized for basic analytic operations, e.g.
 - Consolidation (roll-up)
 - Drill-down (navigate through details)
 - Slicing and dicing (data cubes)

■ We will use the following relation to illustrate OLAP concepts

- *sales (item_name, color, clothes_size, quantity)*

Example *sales* relation

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
dress	dark	small	2
dress	dark	medium	6
dress	dark	large	12
dress	pastel	small	4
dress	pastel	medium	3
dress	pastel	large	3
dress	white	small	2
dress	white	medium	3
dress	white	large	0
pants	dark	small	14
pants	dark	medium	6
pants	dark	large	0
pants	pastel	small	1
pants	pastel	medium	0
pants	pastel	large	1
pants	white	small	3
pants	white	medium	0
pants	white	large	2
shirt	dark	small	2
shirt	dark	medium	6
shirt	dark	large	6
shirt	pastel	small	4
shirt	pastel	medium	1
shirt	pastel	large	2
shirt	white	small	17
shirt	white	medium	1
shirt	white	large	10
skirt	dark	small	2
skirt	dark	medium	5
...
...

Cross Tabulation of sales by item_name and color


clothes_size **all**

		<i>color</i>			
<i>item_name</i>		dark	pastel	white	total
	skirt	8	35	10	53
	dress	20	10	5	35
	shirt	14	7	28	49
	pants	20	2	5	27
	total	62	54	48	164

- The table above is an example of a **cross-tabulation** (**cross-tab**), also referred to as a **pivot-table**.
 - Values for one of the dimension attributes form the *row headers*
 - *Color in the above example*
 - Values for another dimension attribute form the *column headers*
 - Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.

Data Cube

- A **data cube** is a multidimensional generalization of a cross-tab
- Can have n dimensions; we show 3 below

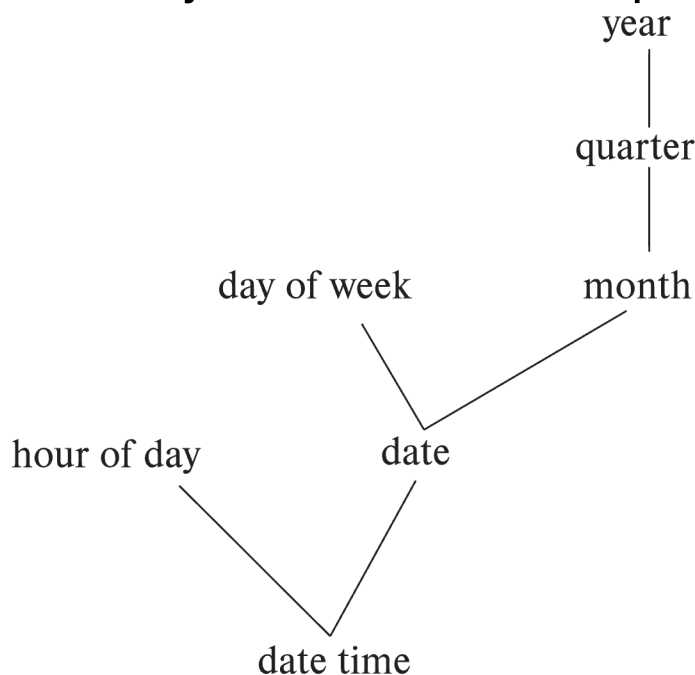


Online Analytical Processing (OLAP) Operations

- **Pivoting:** changing the dimensions used in a cross-tab
 - E.g. moving *colors* to column names
- **Slicing:** creating a cross-tab for fixed values only
 - E.g. fixing *color* to white and *size* to small
 - Sometimes called **dicing**, particularly when values for multiple dimensions are fixed.
 - Goal: divide a quantity of information up into smaller parts, especially in order to analyze it more closely or in different ways.
- **Rollup:** moving from finer-granularity data to a coarser granularity
 - Advanced Aggregation
 - E.g. moving from aggregates by day to aggregates by month or year
- **Drill down:** The opposite operation - that of moving from coarser-granularity data to finer-granularity data

Hierarchies on Dimensions

- **Hierarchy** on dimension attributes: lets dimensions be viewed at different levels of detail
- E.g., the dimension *datetime* can be used to aggregate by hour of day, date, day of week, month, quarter or year



(a) time hierarchy



(b) location hierarchy

Cross Tabulation With Hierarchy

- Cross-tabs can be easily extended to deal with hierarchies
 - Can drill down or roll-up on a hierarchy
 - E.g. hierarchy: *item_name* → *category*

clothes_size:

all

<i>category</i>		<i>item_name</i>		<i>color</i>		
		dark	pastel	white	total	
womenswear	skirt	8	8	10	53	88
	dress	20	20	5	35	
	subtotal	28	28	15		
menswear	pants	14	14	28	49	76
	shirt	20	20	5	27	
	subtotal	34	34	33		
total		62	62	48		164

Relational Representation of Cross-tabs

- Cross-tabs can be represented as relations
- We use the value **all** to represent aggregates.
- The SQL standard actually uses *null* values in place of **all**
 - Works with any data type
 - But can cause confusion with regular null values.

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	white	all	28
shirt	all	all	49
pants	dark	all	20
pants	pastel	all	2
pants	white	all	5
pants	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

OLAP IN SQL

Pivot Operation using SQL

- **select** *
from *sales*
pivot (
 sum(*quantity*)
 for color in ('dark','pastel','white')
)
order by *item name*;
- Not available in MySQL

<i>item_name</i>	<i>clothes_size</i>	<i>dark</i>	<i>pastel</i>	<i>white</i>
dress	small	2	4	2
dress	medium	6	3	3
dress	large	12	3	0
pants	small	14	1	3
pants	medium	6	0	0
pants	large	0	1	2
shirt	small	2	4	17
shirt	medium	6	1	1
shirt	large	6	2	10
skirt	small	2	11	2
skirt	medium	5	9	5
skirt	large	1	15	3

Cube Operation

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- E.g. consider the query

```
select item_name, color, size, sum(number)  
from sales  
group by cube(item_name, color, size)
```

This computes the union of **eight** different groupings of the *sales* relation:

```
{ (item_name, color, size), (item_name, color),  
  (item_name, size),      (color, size),  
  (item_name),          (color),  
  (size),              ( ) }
```

where () denotes an empty **group by** list.

- For each group, the result contains the null value for attributes not present in the group.
- Not available in MySQL

Cube Operation Example

```
select item_name, color,  
sum(number)  
from sales  
group by cube(item_name,  
color)
```

The SQL standard actually uses the **null** value in place of **all**, but to avoid confusion with regular null values, we used **all**

4 groups:

```
{ (item_name, color),  
  (color),  
  (item_name),  
  ( ) }
```

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	white	all	28
shirt	all	all	49
pants	dark	all	20
pants	pastel	all	2
pants	white	all	5
pants	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

Grouping Function

- The function **grouping()** can be applied to an attribute
 - Returns
 - 1, if the value is a **null** value representing **all**
 - 0, in all other cases.
 - Helps to substitute nulls to all
 - **Grouping** (not cube!) works in MySQL
 - More info: <https://dev.mysql.com/blog-archive/mysql-8-0-grouping-function/>
 - Example of use:

```
select case when grouping(item_name) = 1 then 'all'  
         else item_name end as item_name,  
         case when grouping(color) = 1 then 'all'  
         else color end as color,  
         'all' as clothes size,  
         sum(quantity) as quantity  
from sales  
group by cube(item name, color);
```

Grouping Function SQL Flavors

- Instead of *case* function:
 - Can also use *decode()* – in Oracle
 - Can also use *if()* – in MySQL
- in the **select** clause to replace nulls by a value, such as **all**
- E.g., in mysql replace *item_name* in the first query by
if(**grouping(item_name), 'all' , item_name)** as item_name

Rollup Construct: Advanced Aggregation

- **Rollup** extension of the *group by* clause
 - Allows you to include extra rows, e.g. representing the subtotals along with the grand total.
 - You can use a single query to generate multiple grouping sets.
- Generates union on every prefix of specified list of attributes

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name, color, size)
```

Group by rollup generates union of four groups:

```
{ (item_name, color, size), (item_name, color), (item_name), ( ) }
```

- Notes: compare with Cube (was 8!)
 - Order of attributes is important!
 - Rollup is available in MySQL, cube is not
- More info: <https://dev.mysql.com/doc/refman/8.0/en/group-by-modifiers.html>

Rollup Construct: Advanced Aggregation (continued)

- Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- E.g., suppose table *itemcategory*(*item_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)  
from sales, itemcategory  
where sales.item_name = itemcategory.item_name  
group by rollup(category, item_name)
```

would give a summary by (*category*, *item_name*) and by (*category*).

In MySQL syntax of the **rollup** construct is slightly different.

E.g. compare standard SQL , e.g. in SQL Server:

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name, color, size)
```

To MySQL:

```
select item_name, color, size, sum(number)  
from sales  
group by item_name, color, size with rollup
```

Note: You can use *having* and *order by* with *group by ..with rollup*

OLAP Implementations

- OLAP systems that use multidimensional arrays in memory to store data cubes, and are referred to as **multidimensional OLAP (MOLAP)** systems.
- OLAP implementations using only relational database features are called **relational OLAP (ROLAP)** systems
- Hybrid systems, which store some summaries in memory and store the base data and other summaries in a relational database, are called **Hybrid OLAP (HOLAP)** systems.

Reporting and Visualization

- **Reporting tools** help create formatted reports with tabular/graphical representation of data
 - E.g. SQL Server reporting services, MS Power BI, Crystal Reports
- **Data visualization** tools help create interactive visualization of data
 - E.g Tableau, FusionChart, plotly, Datawrapper, Google Charts, etc. etc.
 - Front-end typically based on HTML+JavaScript

**Acme Supply Company, Inc.
Quarterly Sales Report**

Period: Jan. 1 to March 31, 2009

Region	Category	Sales	Subtotal
North	Computer Hardware	1,000,000	1,500,000
	Computer Software	500,000	
	All categories		
South	Computer Hardware	200,000	600,000
	Computer Software	400,000	
	All categories		

Total Sales 2,100,000